

THE NETWORK OF INTENT

How Personal AI Agents Could Rebuild
Discovery, Trust, and Work

AN INVESTMENT THESIS FOR AN AGENT-MEDIATED INTERNET



THE NETWORK OF INTENT

*How Personal AI Agents Could Rebuild
Discovery, Trust, and Work*

An Investment Thesis for an Agent-Mediated Internet

Investment Thesis Edition, 2026

Abstract

The Internet made communication cheap and information abundant, but it did not make it easy for the right people to find one another. We still search manually, perform relevance in public, send cold messages, apply through forms, ask for introductions, attend events, compare uncertain profiles, and pay intermediaries to filter the resulting noise. The same structural problem appears in hiring, fundraising, research collaboration, business development, professional networking, local services, and community formation.

This book proposes a different layer for the Internet: a decentralized network of personal and organizational agents that represent explicit human intentions. These agents would discover one another, exchange questions, verify evidence, negotiate constraints, absorb spam, and bring people into a conversation only when a plausible match exists. The intended product is not a social network with an AI assistant added to it. It is a protocol and market infrastructure for reciprocal discovery.

The opportunity is large, but the thesis is not naive. Agent-mediated networking will create a new arms race in persuasion, identity fraud, synthetic credentials, ranking manipulation, and automated discrimination. A viable system must therefore make real value easier to prove and false claims expensive to sustain. It must use contextual reputation rather than universal scores, selective disclosure rather than mass surveillance, open protocols rather than one central matchmaker, and human approval rather than silent automated exclusion.

The investment thesis is that a useful company can be built before the full network exists. The initial product can operate over today's web, email, professional profiles, research databases, and calendars. Its first beachhead should be a dense, high-value domain in which poor matching is costly and evidence already exists: deep-technology project formation, research consortia, specialized expertise, and early-stage venture collaboration. From there, the network can expand into investor discovery, specialist hiring, and eventually local service markets. The central metric is not engagement. It is the number of qualified human conversations created per unit of attention consumed.

Preface: A Search Problem Disguised as a Social Problem

A founder needs an investor who understands a difficult technology, accepts a long development horizon, and will not force the company into the wrong market. A research group needs two industrial partners, one clinical site, and a coordinator willing to carry administrative responsibility before a European deadline. A company needs a cryptographer who understands production systems rather than only academic primitives. A homeowner needs an electrician who is available this week, licensed for the required work, comfortable with an old building, and trusted by people with similar properties.

These situations look unrelated because they belong to different industries. They use different websites, contracts, vocabularies, and forms of evidence. Yet the underlying task is almost identical. One side has an intention, a set of constraints, and partial information. The other side has capabilities, preferences, availability, and its own uncertainties. Both sides need to discover each other, reveal enough information to test compatibility, negotiate the ambiguous parts, and decide whether a direct conversation is worth the time.

Today, humans perform most of that work manually. We search, scan, compare, write, wait, follow up, translate, verify, and repeat. We often do it while pretending that the process is more rational than it is. A polished profile may stand in for competence. A warm introduction may stand in for fit. Persistence may stand in for seriousness. Visibility may stand in for relevance. A high response rate may stand in for value. The person who can tolerate the most rejection or create the most persuasive presentation can defeat someone whose underlying work is better but whose appetite for self-promotion is lower.

The first promise of the Internet was access. The second was disintermediation. Both were partly fulfilled. Information that once required a library, broker, recruiter, publisher, travel agent, or local directory can now be reached in seconds. Yet the intermediaries did not disappear. Many were replaced by platforms that became larger, more efficient, and more central than the institutions they displaced. They aggregate identity, attention, transactions, and reputation. They are useful precisely because the open Internet does not provide a reliable way to discover the right counterpart or to know whom to trust.

Artificial intelligence creates an opportunity to revisit this unfinished problem. Language models can read unstructured profiles, documents, project descriptions, emails, proposals, reviews, and portfolios. Agents can maintain state, ask questions, use tools, compare alternatives, and interact with other agents. Open standards are emerging for agents to access tools and to communicate across systems. Verifiable credentials and digital identity wallets are becoming practical infrastructure rather than only research topics.^{[11][12][13][14]}

But the obvious implementation is also the dangerous one: place a powerful agent inside an existing centralized platform, give the platform even more data about every user's goals and vulnerabilities, and let it optimize matching according to its own commercial incentives. That would reduce friction while concentrating power. The more useful the agent became, the more the platform would know about jobs we secretly want, investors we fear, projects we cannot yet announce, financial constraints we do not publish, and people we are considering leaving.

This book argues for a different architecture. The agent should represent the person or organization, not the platform. It should be controlled, replaceable, and auditable. Its private memory should not automatically become a corporate dataset. It should be able to speak honestly with its owner and selectively with other agents. It should be able to say no. It should know when it lacks evidence. It should explain why it recommends a conversation and what uncertainty remains. The aim is not to automate human relationships. It is to remove enough discovery and coordination friction that human relationships can begin at a better point.

Executive Investment Thesis

The proposed company would build the first practical layer of an open Network of Intent: software through which individuals and organizations create trusted agents that represent what they seek, what they can offer, what they will not accept, and what evidence they are willing to reveal. The agents search both the existing Internet and the emerging agent network. They engage candidate agents, run reciprocal qualification, exchange selected evidence, negotiate constraints, and request human attention only after a credible basis for interaction has been established.

The core product insight is that matching is not a ranking problem. It is a staged conversation under incomplete information. A job seeker is not merely a vector of skills. An investor is not merely a fund size and sector label. A research partner is not merely a list of publications. A tradesperson is not merely a star rating. Compatibility depends on timing, constraints, incentives, working style, risk, trust, geography, budget, and details that neither side wants to publish globally. The product must therefore combine search with controlled dialogue.

The market insight is that several large platform categories already monetize the same underlying scarcity: trusted access to relevant people. LinkedIn monetizes Talent Solutions, Sales Solutions, Premium subscriptions, and advertising; Microsoft describes demand for these products as a principal driver of LinkedIn revenue.^[5] Recruit Holdings operates Indeed and Glassdoor as a two-sided talent marketplace.^[18] Upwork reported \$787.8 million in 2025 revenue and 785,000 active clients, demonstrating substantial willingness to pay for mediated access to skilled work.^[6] Angi's filings describe a business built around service requests, leads, matching, and booking in home services.^[7] These companies differ, but each captures value because discovery, qualification, and trust remain difficult.

The timing insight is that the required technical components are arriving simultaneously. Agent-to-agent interoperability is moving from isolated prototypes toward open standards; the Linux Foundation reported broad organizational support and production adoption for the A2A protocol in 2026.^[11] The Model Context Protocol provides a standard way for AI applications to connect to tools and data, and its 2026 specification strengthens authorization and extensibility.^[12] W3C Verifiable Credentials 2.0 provides a standardized model for tamper-evident credentials, while the European Digital Identity Wallet architecture is building a large

interoperable ecosystem for storing and presenting attestations.^{[13][14]} None of these standards solves matching, but they reduce the cost of building it.

The initial beachhead should not be mass-market social networking. A new empty social graph has little value, and a general-purpose matching promise is too vague. The first market should have four properties: each successful match is economically meaningful; current search is painful and repeated; there are enough public or verifiable traces to evaluate candidates; and one side can invite the other with a specific opportunity rather than a generic request to join a network.

Deep-tech project formation, research consortia, and specialized venture collaboration satisfy these conditions. A consortium coordinator searching for a cybersecurity SME, a clinical validation partner, or a manufacturing testbed already has a concrete opportunity, a deadline, a budget structure, and role requirements. Potential partners have projects, publications, patents, repositories, certifications, and prior collaborations that agents can analyze. The initial product can search CORDIS, EURAXESS, ORCID, organizational websites, public project databases, documents supplied by users, and existing professional networks. CORDIS itself exposes project and collaboration information, while ORCID is explicitly designed to connect researchers with contributions through persistent identities.^{[8][10]}

The wedge product can be useful with only one customer-side agent. It searches the old Internet, creates a structured shortlist, drafts personalized engagement, handles responses, asks permissioned follow-up questions, and schedules qualified meetings. When a promising external candidate lacks an agent, the system sends an opportunity envelope: a concise private explanation of why the candidate may be relevant, what evidence was used, what remains confidential, and how the candidate can authorize an agent to evaluate the opportunity. The invitation contains value before registration. This is the proposed viral loop.

The company's long-term moat would not be the language model. Foundation models will be substitutable. The defensible assets would be the trusted outcome graph, the domain-specific matching procedures, the permission and policy system, the integrations that let users import evidence without surrendering control, the protocol governance, and a reputation for not selling access or rank. A match network becomes valuable when users believe that it will protect confidential intentions, resist manipulation, and produce better conversations than manual search.

The business model should align revenue with reduced friction rather than increased attention. The recommended early model is subscription

software for organizations and professional teams, with usage tiers based on active searches, agent negotiations, private workspaces, and verified organizational credentials. Managed onboarding and domain-specific matching operations can support early revenue. Later products can include private network deployments, verification services, transaction support, and protocol infrastructure. The company should avoid advertising, paid ranking, and indiscriminate per-lead fees because these recreate the incentives the product exists to remove.

The main risks are severe. A central platform may copy the user experience and exploit an existing graph. Users may not trust agents with private goals. Agents may flood one another with synthetic opportunities. Verification may become expensive and exclusionary. Regulation may treat employment or access decisions as high-risk automated decision-making. An open protocol may commoditize the interface before a durable business is built. The network may optimize for legible evidence and disadvantage unconventional people whose value is difficult to document.

These risks do not invalidate the thesis. They determine the design. The company should begin as an assistive system with human approval, not an autonomous gatekeeper. It should measure qualified conversations and completed collaborations rather than messages sent. It should preserve negative feedback privately and contextually rather than creating permanent public scores. It should allow users to inspect and correct the evidence used about them. It should make reciprocal disclosure a negotiation, not a data extraction event. For employment and similarly consequential domains, it must retain human review and comply with data protection and AI governance requirements.^{[15][16]}

The strongest reason to invest is not that AI can write better outreach. That capability is already becoming common and will increase spam. The stronger thesis is that AI can move the unit of networking from a public profile and a message to a private intention and a structured negotiation. If that shift occurs, the dominant network of the next era may not be the one with the most posts. It may be the one trusted to decide when two people should meet.

Table of Contents

PART I - THE INTERNET'S DISCOVERY FAILURE

- Chapter 1 - Connected, Yet Unable to Find One Another
- Chapter 2 - Networking as Manual Labor
- Chapter 3 - Five Markets, One Structural Problem
- Chapter 4 - Why the Intermediaries Survive

PART II - A NETWORK THAT REPRESENTS US

- Chapter 5 - From Assistants to Representatives
- Chapter 6 - The Mechanics of a Clean Match
- Chapter 7 - Trust Leaves Traces
- Chapter 8 - The Agent Arms Race
- Chapter 9 - Decentralization Without Ideology
- Chapter 10 - The Human Experience of Being Represented

PART III - THE STARTUP THESIS

- Chapter 11 - Choosing the Beachhead
- Chapter 12 - The Minimum Viable Network
- Chapter 13 - Bootstrapping the Network Effect
- Chapter 14 - Business Model, Moat, and Expansion
- Chapter 15 - Risks, Regulation, and Kill Criteria
- Chapter 16 - The Internet That Finds Us
- Selected Sources and Notes

PART I - THE INTERNET'S DISCOVERY FAILURE

Chapter 1 - Connected, Yet Unable to Find One Another

1.1 Distribution Was Solved Before Discovery

The modern Internet is exceptionally good at moving information from one known address to another. It can deliver a document, stream a lecture, publish a profile, transmit a message, and place a video call across continents. It is much less reliable when the address is unknown and the real request is relational: find the particular person, team, organization, or provider that fits a nuanced need.

Search engines reduce this problem to documents. Social networks reduce it to profiles and graph proximity. Marketplaces reduce it to listings, filters, reviews, and bids. Recruiters, consultants, brokers, and community managers add human judgment around these databases. Each mechanism helps. None represents the full intention of both sides.

Consider a founder looking for an investor. The visible query might be: European investor, artificial intelligence, seed stage, checks between EUR 500,000 and EUR 2 million. The real query is longer. The founder wants someone who understands that the product may need four years of technical work before conventional growth metrics become meaningful; accepts a hybrid research and venture model; does not require relocation; can help recruit a commercial leader; is comfortable with open-source components; and has not funded a direct competitor. The founder may also prefer not to reveal publicly that the company has nine months of runway or that an existing cofounder relationship is fragile.

The investor's real query is equally complex. The fund may state that it invests in applied AI, but its internal concerns include partner attention, portfolio concentration, reputational risk, follow-on capacity, conflicts, timing, founder coachability, and whether the proposed market can support the fund's return model. Some constraints are explicit. Others emerge only through conversation. A database can retrieve candidates, but it cannot complete the negotiation.

This distinction explains why discovery remains expensive despite abundant information. The problem is not only that data is missing. Much of the relevant data is conditional. One party will reveal it only if the other

first demonstrates seriousness. Some facts change quickly, such as availability, budget, urgency, and appetite. Some are difficult to encode, such as working style or the kind of disagreement a team can survive. A real match is not a static similarity score. It is a temporary zone of compatibility produced by reciprocal disclosure.

The Internet's dominant interfaces hide this complexity. A search box suggests that the user can describe the answer directly. A profile suggests that a person can be summarized publicly. A five-star rating suggests that trust can be compressed into one number. A recommendation feed suggests that relevance is the same as likelihood of engagement. These simplifications are commercially useful because they create scalable products. They also force humans to perform the missing interpretation outside the system.

That missing work appears in calls, email chains, introductions, proposal documents, interview loops, reference checks, test assignments, background research, and informal conversations. It is where fit is actually discovered. It is also where time is lost.

1.2 The Hidden Coordination Tax

The cost of poor discovery is rarely recorded as one line in a budget. It is distributed across many people and appears as normal work. A founder spends two hours identifying investors, three hours researching their portfolios, two hours finding contact paths, and another hour adapting messages. An investor reads dozens of nearly identical pitches to find one that deserves a second look. A research coordinator asks partners for descriptions, discovers late that one cannot satisfy an eligibility rule, and rebuilds the consortium. A recruiter screens hundreds of applications that were easy to submit but weakly related to the role. A homeowner requests five quotes and receives two, neither comparable.

Because each individual task is small, the aggregate tax is easy to ignore. Yet it affects almost every collaborative activity. The tax has at least six components.

The first is search cost: finding a plausible set of candidates across fragmented sources. The second is interpretation cost: understanding whether labels and claims mean the same thing in context. The third is verification cost: determining whether the person, opportunity, credential, review, or organization is real. The fourth is engagement cost: writing messages, requesting introductions, and overcoming inbox defenses. The


fifth is negotiation cost: revealing constraints and testing alignment. The sixth is opportunity cost: the valuable work not performed while all the previous steps consume attention.

Platforms often optimize one component while worsening another. Easy application tools reduce the applicant's submission cost, but they can increase screening volume for employers. Automated outreach reduces the sender's engagement cost, but it raises verification and filtering costs for recipients. Lead-generation marketplaces simplify discovery for customers, but may sell the same request to multiple providers, producing a race to contact the customer first. Public recommendation systems improve browsing, but reward actors who optimize visibility rather than actors who quietly deliver value.

LinkedIn's own design illustrates the tension. Job seekers can use Easy Apply to submit without leaving the platform, while recruiters use advanced search filters and messaging tools to find candidates.^{[1][2]} Sales Navigator similarly emphasizes finding and engaging the right leads through extensive filters.^[3] These are rational responses to the same market need. Yet LinkedIn introduced limits on Easy Apply submissions specifically to encourage more intentional applications and curb automation and bots.^[4] Lowering friction on one side had created too much low-cost activity for the system to absorb.

This is not a criticism of one company. It is a general law of networked markets: when sending is cheaper than evaluating, volume grows until evaluation becomes the scarce resource. Generative AI accelerates that law. It can produce personalized messages, tailored resumes, project summaries, proposals, and follow-ups at negligible marginal cost. If recipients still evaluate manually, the result is not a frictionless market. It is an arms race in plausible text.

A Network of Intent must therefore optimize the cost of the pair, not merely the convenience of the sender. A successful system should make it cheap to express a genuine intention while making it costly to create ungrounded claims, irrelevant requests, or deceptive identities. It should reduce total attention consumed per successful collaboration.



The current discovery funnel concentrates human effort before either side knows whether a meaningful match exists.

1.3 What a Good Match Actually Contains

A match is often described as compatibility between attributes. The candidate has a skill; the project requires it. The investor has a thesis; the startup belongs to the category. The plumber serves the postal code; the customer lives there. This is necessary but insufficient.

A useful match has at least five layers. The first is capability: can the party do what is required? The second is intention: does the party actually want this type of opportunity now? The third is constraint compatibility: do budget, timing, geography, legal status, risk, and working arrangements fit? The fourth is evidence: is there enough reason to believe the claims? The fifth is relational viability: can the parties communicate, negotiate, and coordinate effectively?

Current systems handle the first layer reasonably well when attributes are standardized. They handle intention poorly because profiles are persistent while intentions are temporary. They handle constraints through forms that often appear too early or too late. They handle evidence through reviews, credentials, portfolios, and references, but each signal can be gamed or interpreted out of context. They handle relational viability mainly by forcing people into meetings.

The proposed network treats these layers as stages. It does not reveal everything at once. It asks the minimum question required to decide whether the next exchange is justified. An agent might first determine that a researcher has relevant expertise and is allowed to participate in the target

funding program. Only then does it ask whether the researcher is available during the proposal period. Only after a positive answer does it request a selected publication, a prior project role, or a short explanation of the proposed contribution. Human contact begins after enough uncertainty has been reduced to make the conversation potentially valuable.

This sequence matters because privacy and efficiency are related. Broadcasting more information can make search easier, but it exposes users and creates new opportunities for manipulation. A well-designed agent network should achieve useful discovery with less public disclosure, not more. It should let a person say privately, "I would consider leaving my job for a role with these characteristics," without announcing that intention to an employer or an advertising platform. It should let an organization explore a partnership without publishing strategy. It should let an investor decline a deal without training a public reputation model on confidential reasoning.

The deeper opportunity is to make intentions first-class objects on the Internet. Documents, profiles, products, and posts already have standard representations. Intentions do not. They are expressed indirectly through search queries, job listings, requests for proposals, messages, posts, wish lists, and forms. They are often copied into multiple systems and become stale immediately.

An intent object would be private by default, time-bounded, revocable, and structured enough for agents to reason about. It could contain a goal, hard constraints, soft preferences, acceptable evidence, disclosure policy, authority limits, expiry, and a definition of success. It could evolve through conversation. The owner could inspect and edit it. Different agents could compete to represent it.

Once such objects exist, the Internet can begin to search not only what has been published, but what people are prepared to do.

Chapter 2 - Networking as Manual Labor

2.1 LinkedIn and the Performance of Relevance

Professional networking on the contemporary Internet is organized around a public profile, a social graph, a stream of content, and paid tools for search and contact. LinkedIn is the clearest example because it serves several markets at once: identity, recruiting, sales, career discovery, advertising, learning, and professional publishing.

The profile performs two functions that are easily confused. It records parts of a person's history, and it acts as a persuasive page designed for unknown future readers. Users choose titles, summaries, featured items, endorsements, recommendations, and visible activities not only to describe themselves but to be discovered and judged. The platform encourages completeness because more structured data improves search and monetization. Users learn to write for recruiters, clients, investors, collaborators, and algorithms simultaneously.

This creates a performance layer around professional life. A person may be excellent at building distributed systems but poor at packaging that work into recognizable keywords. Another may understand exactly how to appear in searches for "AI strategy" despite having little evidence of execution. A founder may post confident progress while privately changing direction. A scientist may have substantial influence through mentoring, datasets, and technical review that is barely visible in publication lists. The public profile is neither false nor complete. It is a negotiated interface between identity and opportunity.

Discovery then depends on the vocabulary of the platform. Recruiters search titles, skills, geography, seniority, companies, schools, keywords, and network relationships. LinkedIn Recruiter advertises more than forty advanced filters, and Sales Navigator offers a similar logic for commercial leads.^{[2][3]} These tools are powerful because they convert a vast population into a queryable database. They also privilege what can be standardized and publicly stated.

After search comes engagement. A connection request, comment, post interaction, warm path, or InMail must cross the recipient's attention

boundary. The message is expected to be brief, specific, personalized, and confident. At scale, senders use templates and automation. Recipients become suspicious of phrases that once signaled care. Personalization becomes another generated surface.

The platform can add AI to each step: rewrite the profile, suggest candidates, generate messages, summarize applicants, predict interest, and recommend introductions. This will make the current system faster. It does not necessarily change its structure. The user still performs publicly for a central database. The sender still pushes a message into another person's queue. The platform still benefits from more activity, more profile data, and continued dependence on its graph.

The alternative proposed here is not to abolish profiles or social graphs. They are useful evidence sources. It is to subordinate them to private intentions. A user's agent could import selected profile data, correct it, add non-public constraints, and search across multiple sources. It could represent the user without requiring the user to post continuously or optimize every sentence for visibility. The public profile would become one input among many rather than the primary economic identity.

This difference is subtle but fundamental. In the current model, people adapt themselves to the network's representation. In the proposed model, the network adapts its search to the person's intention.

2.2 Email and the Cold Outreach Economy

Email remains the universal fallback because it is open enough to reach almost anyone and mature enough to cross organizational boundaries. It is also where the failure of discovery becomes most visible.

A typical outreach workflow begins outside email. The sender identifies a category of people, searches databases and websites, collects addresses, verifies them, segments the list, researches recipients, drafts a message, and sequences follow-ups. Specialized products exist for every step. The sender measures opens, replies, meetings, and conversions. The recipient sees only the final message and must decide whether it is legitimate, relevant, safe, and worth answering.

The asymmetry is severe. One sender can generate thousands of plausible messages. Each recipient must evaluate them individually. Spam filters can block obvious abuse, but high-quality automated outreach is designed to resemble genuine human communication. The better generative

systems become, the less linguistic polish indicates that a sender invested attention.

Cold email is not inherently bad. Many valuable relationships begin without an introduction. The problem is that the protocol provides almost no structured way to establish relevance before the message enters the inbox. The sender cannot easily prove why the recipient was selected. The recipient cannot ask an agent to negotiate without exposing an address or replying. There is no standard mechanism for reciprocal disclosure, evidence requirements, or rate-limited access to attention.

As a result, people develop crude defenses. They ignore unknown senders, delegate inbox review, publish contact forms, hide addresses, rely on assistants, or require warm introductions. These defenses protect attention but reduce openness. Someone outside the right network may have a genuinely important reason to make contact and still be filtered out.

An agent network could preserve the openness of email while adding a pre-inbox qualification layer. The sender's agent would create a signed intent message containing a minimal claim: the purpose of contact, why the recipient appears relevant, what evidence supports that inference, what is being requested, and what the sender is willing to disclose. The recipient's agent could accept, reject, request clarification, impose a challenge, or route the request to another representative. The human inbox would receive only those exchanges that pass the user's policy.

During the transition, this can work through ordinary email. The agent may send a readable message with a machine-readable attachment or link. The recipient does not need to join a new network to understand the request. If the recipient has a compatible agent, the agents continue directly. If not, the message still functions as a higher-quality introduction.

This transitional capability is strategically important. A startup cannot wait for universal adoption. It must use existing channels while gradually replacing their manual coordination burden.

2.3 Warm Introductions, Events, and Human Brokers

Warm introductions persist because they compress trust. When a respected person connects two others, the recipient infers that the sender is real, that some relevance check occurred, and that ignoring the message may carry a social cost. The introducer lends reputation to the interaction.

This mechanism is effective but scarce. It favors people already embedded in influential networks. It also consumes the introducer's attention and creates informal obligations. Founders repeatedly ask the same investors for introductions. Researchers depend on senior colleagues to enter consortia. Candidates seek referrals inside companies. Service providers rely on local word of mouth. The result is a social routing system that works well within clusters and poorly across their boundaries.

Conferences, meetups, demo days, trade shows, and professional associations address the same problem by creating temporary high-density environments. Participants pay with travel, time, sponsorship, and attention. They browse badges, schedules, booths, speaker lists, and side events. The most valuable conversations are often serendipitous, but serendipity is expensive when thousands of people are present and each has only a few hours.

Professional brokers formalize the routing function. Recruiters, executive search firms, investment bankers, talent agents, consortium consultants, real-estate brokers, and marketplace operators maintain networks and interpret needs. Good brokers do more than supply names. They translate between parties, understand hidden constraints, protect confidentiality, and manage timing. Their value demonstrates that matching is conversational and contextual.

A credible agent network should learn from brokers rather than dismiss them. The aim is not to claim that software can instantly replace human judgment. It is to encode the repetitive parts of brokerage and let human experts intervene where ambiguity, accountability, or relational skill matters most.

In fact, brokers may become important early customers and ecosystem participants. A recruiter could operate an agent with domain knowledge and a trusted candidate network. A research consultant could publish a capability agent that understands program rules and consortium structure. A local trade association could certify members and route requests. The open network would allow specialized intermediaries to compete on judgment and trust without owning the user's identity or locking the relationship inside one marketplace.

The likely future is not disintermediation in the literal sense. It is a change in what intermediaries are paid to do. Low-value list building, repeated screening, and status chasing can be automated. High-value interpretation, accountability, conflict resolution, and domain expertise remain valuable. The intermediary becomes an agent operator, verifier,

guarantor, or exception handler rather than a gatekeeper to a proprietary directory.

Chapter 3 - Five Markets, One Structural Problem

3.1 Applying for Jobs

Job search is one of the clearest examples of a market damaged by asymmetric automation. Employers publish descriptions assembled from organizational needs, legacy templates, compliance language, and aspirational requirements. Candidates search by title, location, salary, keywords, and company. They tailor resumes, answer forms, write cover letters, complete assessments, and wait. Employers use applicant tracking systems, screening questions, recruiter searches, referrals, and interviews to reduce the resulting volume.

Easy Apply illustrates both the benefit and the limit of lower friction. LinkedIn lets candidates submit directly through the platform, while some jobs redirect to an employer site.^[1] The convenience is real. It also makes broad application strategies cheaper. LinkedIn's later limits on Easy Apply submissions acknowledge the need to restrain bots and encourage more intentional behavior.^[4]

The obvious AI product automates applications. It identifies jobs, adapts resumes, writes answers, and submits at scale. That helps one user temporarily, but it makes the market worse if every candidate does it. Employers respond with more screening, stronger filters, automated interviews, identity checks, and assessments. Candidates then use AI to prepare for those filters. The process becomes a contest between application agents and screening agents, while both sides still know little about actual fit.

A reciprocal system would begin differently. The candidate creates a private work intent: desired problems, acceptable roles, salary range, location constraints, availability, working style, learning goals, exclusions, and evidence. The employer creates a role intent that separates hard constraints from preferences and distinguishes the work actually required from conventional labels. Agents test compatibility before an application is created.

The candidate may disclose salary expectations only after the employer confirms the range. The employer may reveal team composition only after

the candidate demonstrates relevant experience. The candidate agent can ask whether a nominally remote role requires unannounced travel. The employer agent can request a work sample that is proportionate to the uncertainty. Both agents can identify contradictions: a company advertising autonomy while requiring daily approval, or a candidate seeking research freedom while preferring narrowly specified tasks.

The human interview then begins with the unresolved questions rather than repeating the resume. This does not remove judgment. It improves the input to judgment.

Employment is also the domain where design errors can cause the most direct harm. Agents could reproduce discrimination, infer sensitive traits, hide opportunities, or create opaque rejections. The European AI Act treats certain employment uses as high-risk, with specific timelines and obligations, while GDPR Article 22 protects people against certain solely automated decisions with significant effects.^{[15][16]} A startup should therefore enter employment after it has developed strong explanation, audit, consent, and human-review mechanisms in less consequential domains.

3.2 Searching for Investors

Fundraising combines database search, social hierarchy, narrative performance, and timing. Founders use investor websites, portfolio pages, databases, accelerator networks, LinkedIn, conference lists, and introductions. They classify funds by stage, geography, sector, check size, lead behavior, and reputation. They prepare decks, memos, data rooms, metrics, and carefully timed outreach.

The visible process is inefficient because investor labels are broad and behavior changes with market conditions. A fund may say it invests in deep tech but prefer projects with near-term commercial proof. A corporate venture arm may appear strategically aligned but move too slowly for the startup's runway. An angel may have relevant operating experience but no appetite for follow-on investment. A founder may say they seek capital while actually needing a chairperson, commercial cofounder, customer access, or patient governance.

Warm introductions dominate because they signal legitimacy and reduce the investor's search burden. They also reinforce network concentration. Founders outside established hubs can spend months trying to reach people whose mandate was never compatible.

An investor-matching agent should not begin by generating more pitch emails. It should construct a mutual fit model. On the founder side, this includes the real financing need, acceptable governance, strategic dependencies, time horizon, desired help, disclosure boundaries, and reasons the company might reject capital. On the investor side, it includes portfolio construction, available reserves, partner expertise, current pace, conflicts, return horizon, decision process, and willingness to support the specific risk profile.

The agent can then stage disclosure. A startup may initially share a redacted technical and market summary. An investor agent can explain what evidence would justify further access. The founder agent can test whether the fund's behavior in prior investments matches its stated philosophy, using public traces and references. Once both sides pass a threshold, the agents open a data-room workflow and schedule a meeting with a clear agenda.

This system would reduce the importance of pitch theater, but it would not eliminate narrative. Investors legitimately need to understand how founders think, persuade, recruit, and respond to uncertainty. The point is not to replace conviction with a score. It is to ensure that conviction is evaluated after basic incompatibilities have been removed and evidence has been organized.

The largest product danger in this market is pay-to-play ranking. Founders are vulnerable, and investors have scarce attention. A platform can easily charge startups for prominence or sell lead access. Doing so would corrupt the trust layer. A Network of Intent must make it explicit that payment buys representation and workflow, not a better place in another party's shortlist.

3.3 Building Research and Innovation Consortia

Research collaboration is often presented as an intellectual process, but consortium formation is also a logistical and political one. A funding call defines expected outcomes, eligibility rules, budgets, timelines, technology readiness, impact requirements, geographic constraints, and reporting obligations. A coordinator must translate these into roles and find organizations willing and able to fill them.

The search happens through previous projects, personal networks, conferences, national contact points, CORDIS, EURAXESS, mailing lists, consultants, and direct outreach. CORDIS provides a public repository of EU-

funded projects, participants, results, and collaboration networks, while EURAXESS offers jobs, funding, hosting, and partnering tools.^{[8][9]} ORCID helps identify researchers and connect them to contributions.^[10] The information exists, but it is fragmented and rarely expresses current intention.

A university may have relevant publications but no administrative capacity to join another proposal. An SME may have the technology but need funded development rather than unpaid proposal work. A hospital may be scientifically ideal but unable to complete ethics approvals in time. A previous project participant may look experienced while the people responsible have left. A consortium may appear complete and still lack a partner willing to own a difficult work package.

An agent-mediated workflow can represent the call itself as an intent object. The coordinator defines required capabilities, eligible countries, role boundaries, budget assumptions, evidence, deadlines, and open uncertainties. The agent searches public sources and private networks, then contacts organizational agents with a precise role hypothesis. Candidate agents respond with availability, constraints, relevant evidence, and questions. The system identifies missing coverage, duplicated capability, dependency risks, and likely negotiation conflicts.

Instead of twenty vague introductory calls, the coordinator receives five candidates whose agents have already confirmed interest, eligibility, timing, and a plausible contribution. The first human meeting can focus on scientific and strategic fit.

This is a strong beachhead because the value of a correct match is high, the cost of current search is visible, and successful invitations naturally recruit new participants. It is also imperfect. The market is cyclical around calls, procurement can be complex, and many research organizations have slow decision processes. The startup must therefore design for reusable organizational agents that support multiple opportunities, not one-off proposal consulting.

3.4 Finding Skilled Trades and Local Services

At first glance, finding an electrician has little in common with forming a research consortium. The vocabulary and transaction size differ, but the matching stages are similar.

A homeowner describes a problem with incomplete technical knowledge. Search engines and directories produce providers. Marketplaces collect service requests and route them to professionals. Reviews, licenses,

photos, response time, location, and price help the customer choose. Providers pay through advertising, subscriptions, commissions, or lead fees. Angi's public filings describe service requests, leads, matching, booking, and independently established providers as core parts of its model.^[7]

The friction is not only discovery. The customer may not know what work is required. The provider needs photos, building details, access conditions, urgency, materials, and whether another trade must act first. Quotes are difficult to compare because scope differs. Availability changes quickly. Reviews may reflect punctuality rather than technical quality, or a provider may be excellent for new construction and unsuitable for historical buildings.

A household agent could maintain private information about the property: age, systems, prior repairs, warranties, preferred materials, access rules, and budget policy. A tradesperson's agent could maintain qualifications, service radius, equipment, current availability, job preferences, and evidence from completed work. The agents could clarify the problem, request safe diagnostic information, determine whether an emergency service is needed, compare scope, and schedule an inspection.

Trust must remain local and contextual. A universal score is not enough. The relevant evidence may include a license, insurance, verified completion of similar work, supplier relationships, and attestations from nearby customers with comparable properties. A new provider should be able to enter through stronger verification or supervised jobs rather than being permanently excluded by lack of historical reviews.

This market is attractive because people repeatedly experience the pain. It is also operationally demanding. Safety, liability, payments, local regulation, disputes, and physical-world verification are difficult. It should therefore be an expansion market after the network has mature identity, evidence, scheduling, and dispute mechanisms.

3.5 Beyond Transactions

Not every useful connection is a job, investment, contract, or purchase. People seek mentors, collaborators, volunteers, study partners, patients with similar rare conditions, local mutual-aid groups, artistic collaborators, peer reviewers, advisors, and communities where a difficult interest is understood.

Current platforms often force these needs into content feeds. A person posts publicly, joins groups, follows hashtags, searches forums, or asks

friends. Discovery depends on visibility and luck. The platforms optimize for activity because activity can be monetized, even when the user's real goal is one durable relationship.

An intent network could support non-commercial goals without turning them into a marketplace. A retired engineer might offer two hours a month to mentor teams working on energy systems. A doctoral student may seek a reading partner for a narrow field. A patient may want to speak with someone who has navigated a particular treatment decision, under strict privacy rules. A community may need three volunteers with specific skills for one weekend.

The system should not infer that every available capability must be monetized. Users need explicit control over which intentions are commercial, reciprocal, civic, private, or exploratory. This is one reason the infrastructure should not be funded primarily through advertising or transaction extraction. A business model that requires every match to generate a fee will neglect or distort some of the network's most human uses.

The investment thesis begins with economically valuable professional coordination because that can fund the infrastructure. The larger purpose is to make the Internet better at forming useful relationships, including those that markets do not price well.

Chapter 4 - Why the Intermediaries Survive

4.1 They Aggregate Scarcity, Not Just Data

It is easy to criticize intermediaries as rent collectors. It is more useful to ask what scarce resource they aggregate.

LinkedIn aggregates professional identity, graph context, and attention. Job boards aggregate open roles and candidate traffic. Investor databases aggregate firm information and deal histories. Recruiters aggregate trust, interpretation, and access. Research portals aggregate calls, projects, and institutional records. Home-service marketplaces aggregate consumer demand, local providers, reviews, and booking. Conferences aggregate people in time and place.

The underlying scarcity is not information alone. It is credible availability. A directory may list ten thousand experts, but only a few are relevant, interested, free, and safe to contact. A platform creates value when it turns an unbounded search into a manageable set of possible actions.

This explains why simple decentralization often fails. Publishing data openly does not automatically create a usable market. Someone must normalize it, update it, rank it, protect it, resolve abuse, and provide a reliable interface. A protocol without product quality becomes an empty standard. A product without governance becomes another central gatekeeper.

The proposed startup must therefore perform real intermediary functions from the beginning. It must route requests, enforce policies, verify selected evidence, manage consent, and provide recourse. The claim is not that intermediaries disappear. The claim is that users should be able to change intermediaries without losing identity, history, and access to the network.

4.2 They Monetize Friction

Many platform business models depend on the friction they reduce. Recruiters are paid because hiring is difficult. Investor databases are valuable because knowledge is fragmented. Marketplaces charge because

trust and transaction management are scarce. Lead-generation systems earn money by controlling access to demand. Social networks monetize visibility because organic reach is uncertain.

This creates an incentive problem. A platform may improve matching while preserving enough opacity that users continue to pay for filters, boosts, messages, and access. It may optimize the amount of activity rather than the speed with which users leave the platform having solved their problem. A dating service, job board, or professional network can face a structural tension: the user's success may reduce future engagement.

The Network of Intent should make a different promise. Its unit of value is a resolved intention. A search can end with a match, a justified rejection, a revised requirement, or evidence that no suitable counterpart is currently available. Each outcome saves attention. The product should be willing to say that the market does not contain a credible match rather than manufacture a list to appear useful.

This requires metrics that differ from conventional platform metrics. Daily active users, time on site, messages sent, and feed impressions are poor primary indicators. Better metrics include time to qualified conversation, percentage of agent engagements that reach reciprocal disclosure, human acceptance of shortlists, meeting-to-collaboration conversion, reduction in repeated questions, and the rate at which users correct agent assumptions.

A company funded by subscriptions can align with these metrics more easily than one funded by ads or paid placement. Subscription is not automatically ethical, but it makes the customer relationship clearer. Organizations pay because the agent reduces coordination cost and improves decisions. Individuals may receive a basic agent as infrastructure, with professional features paid by teams or counterpart organizations.

4.3 Why AI Inside Existing Platforms Is Not Enough

Existing platforms will add increasingly capable agents. Recruiters will ask a hiring assistant to find candidates. Sales teams will ask agents to identify accounts and draft outreach. Job seekers will ask agents to search and apply. Marketplaces will summarize requests and recommend providers. These products will be useful and will validate the broader behavior.

The investment case for a new network depends on what centralized platform agents cannot credibly provide.

First, a platform agent is bounded by the platform's data and incentives. It can integrate external sources, but the business benefits when the user remains dependent on its graph. Second, it cannot easily represent both sides without conflicts. An agent paid by an employer and hosted by a recruiting platform may claim to help candidates, but its primary optimization is not neutral. Third, users may hesitate to reveal private intentions to a platform whose advertising, ranking, or enterprise products depend on those intentions. Fourth, portability is weak. A user who changes provider may lose history, trust signals, and learned preferences.

An open intent network can differentiate through fiduciary-like representation: the agent's policy is explicitly tied to the user. The user can export the intent model, credentials, negotiation history, and preferences. Other providers can compete to operate the agent. The protocol separates identity and communication from the application interface.

This position is not easy to defend. A large platform can offer convenience and immediate graph density. The startup must therefore interoperate with existing platforms before it challenges them. It should let users connect LinkedIn, email, calendars, repositories, CRM systems, research identifiers, and documents. It should produce value from those sources while keeping a private, portable layer above them.

The strategic analogy is not a social network replacing another social network. It is email operating across providers, the web operating across publishers, or payment protocols operating across banks and applications. The company can be important without owning every node, provided it owns a trusted product, a strong implementation, and influential protocol governance.

PART II - A NETWORK THAT REPRESENTS US

Chapter 5 - From Assistants to Representatives

5.1 The Difference Between Helping and Representing

Most personal AI products are described as assistants. They answer questions, summarize documents, draft messages, search calendars, and execute tasks. The word is comfortable because it implies a subordinate tool. A networking agent needs a stronger concept. It must represent its owner in interactions where the other side has its own interests.

Representation changes the design in four ways.

First, the agent needs authority boundaries. It must know what it may disclose, promise, schedule, decline, or negotiate. A drafting assistant can produce a sentence and wait. A representative may need to refuse a request before the human sees it. That power must be explicit, revocable, and narrow.

Second, the agent needs a durable model of interests rather than only a conversation history. It should understand hard constraints, soft preferences, unresolved tensions, and priorities that change by context. A founder may accept a lower valuation from a strategically valuable investor but reject the same terms from a passive fund. A researcher may join a project only if it funds a specific line of work. A homeowner may prefer the cheapest repair for a rental property and the most durable repair for a family home.

Third, the agent must distinguish persuasion from truth. Its task is not to maximize the chance that every request succeeds. It should help the owner present relevant evidence clearly, but it should not invent fit. A trustworthy representative sometimes concludes that the owner is not suitable, that the opportunity is misdescribed, or that more preparation is required.

Fourth, the agent must be accountable to the owner. The user should be able to inspect why a person was contacted, what information was shared, which claims were made, and what policy was applied. A human representative can be questioned. A digital one should leave a better audit trail.

This suggests a design principle: the agent should optimize long-term represented interest, not local conversion. It may choose not to send a message that could win a meeting but damage reputation. It may protect the user from a prestigious opportunity whose constraints are incompatible. It may recommend an honest admission of uncertainty because the evidence is weak.

The distinction is commercially important. A generic assistant will become a feature of operating systems, email clients, CRMs, and professional platforms. A representative is a product relationship built on trust, policy, and accumulated understanding. Users may switch models while retaining the representative layer.

5.2 The Intent Model

The central data structure of the proposed network is an intent model. It is not a profile and not merely a prompt. It is a versioned description of a goal and the conditions under which the agent may pursue it.

A practical intent model contains at least ten elements.

The goal states the desired change: find an investor, recruit a specialist, form a consortium, hire a service provider, locate a mentor, or explore a collaboration. The offered value states what the owner contributes. Hard constraints define conditions that cannot be violated. Soft preferences express trade-offs. Evidence references claims that the agent may use. Disclosure rules determine who may see what and at which stage. Authority rules define actions the agent may take. Time bounds prevent stale intentions from circulating indefinitely. Risk limits constrain commitments, identity exposure, and external actions. Success criteria define when the intent can close.

These elements need not appear as a complex form. The user can describe the situation in ordinary language. The agent asks clarifying questions and converts the answers into a structured model. The user then reviews a short representation of the model, especially the hard constraints and disclosure policy.

For example, a founder might say:

We are looking for EUR 1.5 million to complete an industrial pilot. I prefer an active investor who can help recruit a CEO, but I do not want a forced relocation or a three-year exit assumption. We can disclose our technical memo after the investor confirms check size, relevant

portfolio experience, and no conflict. Do not contact direct competitors or funds that require a consumer growth model.

The agent extracts financial range, stage, use of funds, desired contribution, governance constraints, time horizon, conflict policy, disclosure sequence, and excluded categories. It should also identify ambiguity. What does "active investor" mean? How much control is acceptable? Is a strategic corporate investor acceptable? What counts as a direct competitor? The purpose of structuring is not to pretend the ambiguity does not exist. It is to make the ambiguity visible early.

Intent models can be layered. A public layer might reveal that an organization seeks a cybersecurity partner for a European research proposal. A network layer might reveal the target call, role, budget range, and country constraints to authenticated participants. A bilateral layer might reveal the draft work package and internal concerns only after reciprocal qualification. A human-only layer might contain political or relational considerations that the agent uses but never discloses.

This layered structure lets the system search widely without broadcasting everything. It also supports multiple agents. A user may authorize one agent to search public sources, another to negotiate private terms, and a human advisor to approve final contact. The system can preserve separation of duties for sensitive decisions.

5.3 Honest Private Language

The quality of representation depends on whether users can speak honestly to their agents. Public professional language is distorted by audience. People say they are "open to opportunities" when they are actively trying to leave. Companies say they seek "passionate self-starters" when they need someone willing to repair a disorganized system. Investors say they are founder-friendly while privately prioritizing ownership and return timing. Service providers say no job is too small while avoiding low-margin visits.

A private agent should create a space for less polished language. The user can say: "I am tired of managing this company and need someone who can take over without destroying its culture." Or: "I need a job within four months, but I cannot sustain another chaotic startup." Or: "I do not trust the contractor's estimate and need a second opinion before I confront them."

This honesty is not a license for the agent to reveal vulnerabilities. It is information used to protect the user. The agent translates internal language

into external criteria and questions. It may say to a candidate CEO: "The founders seek an operator with genuine authority and a gradual transition plan," not "the founder is exhausted." It may ask an employer about decision processes and turnover without disclosing the candidate's previous burnout.

The system therefore needs two semantic layers: the owner's private model and the negotiable external representation. A simple chatbot history is insufficient because accidental disclosure would be catastrophic. Policies must be machine-enforceable, and sensitive facts should be tagged by purpose, recipient class, expiry, and required approval.

A useful product feature would be a disclosure rehearsal. Before an agent enters a new search, it shows the user examples of what it would say at each stage. The user can tighten or relax the policy. This makes representation concrete and reveals misunderstandings before they reach another party.

The long-term competitive advantage may depend less on how eloquently the agent speaks than on how reliably it keeps these boundaries.

Chapter 6 - The Mechanics of a Clean Match

6.1 Matching as a Protocol, Not a Feed

The user experience can be simple even if the underlying protocol is not. A person should be able to state a goal, inspect the interpreted intent, approve a search, and receive a small number of justified candidates. Behind that interface, the agents conduct a staged process.

A clean match can be modeled in six phases: private intent, discovery, reciprocal qualification, evidence exchange, constraint negotiation, and human handoff.

In the private intent phase, each side defines goals, policies, and authority. In discovery, an agent searches public data, private networks, federated indexes, and existing platforms using only information allowed for broad routing. Reciprocal qualification begins when one agent contacts another with a minimal, signed explanation of relevance. The receiving agent checks the request against its owner's active intentions and policies. If there is possible alignment, the agents exchange questions. Evidence exchange follows only when claims matter to the decision. Negotiation tests timing, terms, responsibilities, and unresolved constraints. The human handoff occurs when both agents can explain why a direct conversation is justified.

A CLEAN AGENT-MEDIATED MATCH



A clean match is a staged protocol. Each phase earns the right to request more information and more human attention.

The sequence should be interruptible. Either side can decline without exposing unnecessary details. The agents can record a reason privately or share a coarse reason such as timing, geography, conflict, or insufficient evidence. A rejection is a legitimate resolution, not a product failure.

This matters because current matching products often produce lists even when the underlying fit is weak. A list demonstrates activity. A protocol can demonstrate restraint. The system earns trust by showing that it can stop.

6.2 Discovery Without Immediate Disclosure

Discovery needs enough information to route an intention, but not enough to reveal the owner unnecessarily. The system can use a combination of public evidence, user-authorized sources, and privacy-preserving query techniques.

At the simplest level, the initiating agent searches public data. For a research partner, it examines projects, publications, websites, repositories, and declared capabilities. For an investor, it examines portfolio, stage, geography, public statements, and known transactions. For a tradesperson, it examines licenses, service areas, categories, and verified work history.

The agent then creates a candidate hypothesis: why this counterpart may fit, what evidence supports the hypothesis, and what must be confirmed. It does not send a generic message. It sends a request grounded in that hypothesis.

When both sides are already in the network, discovery can be more private. Agents can publish coarse capability cards or encrypted routing descriptors that reveal categories without exposing detailed intentions. A research organization might advertise that it is open to selected applied-AI projects in health and manufacturing during a given quarter, without publishing specific strategic interests. A candidate might permit searches for senior technical roles in two countries while keeping the job search invisible to public networks.

Future implementations could use private set intersection, secure computation, or cryptographic commitments for narrow checks, but the startup should not make advanced cryptography a prerequisite for the first product. Most initial privacy gains come from disciplined data minimization, local storage, staged disclosure, and clear policy.

Cryptographic techniques should be added where they solve a concrete threat model.

6.3 Reciprocal Qualification

Traditional forms are unilateral. The employer asks; the candidate answers. The investor requests a deck; the founder supplies it. The marketplace asks the homeowner to describe the job; providers compete. Reciprocal qualification recognizes that both sides need assurance.

An agent conversation can interleave questions. A project coordinator asks whether an SME has relevant experience and capacity. The SME asks whether the budget covers integration work and whether the coordinator has submitted similar proposals. An investor asks about traction and technical risk. The founder asks about decision timing, reserves, and behavior in difficult portfolio situations. A homeowner asks for proof of insurance. The provider asks whether access is possible and whether asbestos may be present.

The protocol should support question provenance. Each question is linked to a decision need: eligibility, risk, scope, preference, or evidence. The receiving side can see why the information is requested and can answer, refuse, or propose a less intrusive alternative.

This prevents qualification from becoming extraction. A party should not have to provide a full data room merely to learn that the counterpart cannot meet a basic constraint. The protocol can enforce symmetry or proportionality: deeper access requires deeper reciprocal disclosure.

The system can also detect repeated questions and reuse verified answers with permission. A company should not ask a candidate to re-enter the same employment history into five systems. A researcher should not repeatedly type institutional data already verified by an identifier or prior project. Reuse saves time, but the owner must remain able to correct stale information and decide who receives it.

6.4 Evidence Exchange and Trials

Claims become useful when they can be connected to evidence. Evidence does not need to be cryptographic in every case. It can include public artifacts, references, verified credentials, signed attestations, work samples, transaction records, licenses, publications, code, project deliverables, or short trials.

The agent should request evidence only when a claim affects the decision. If a partner says it has coordinated a European project, a public project record may be sufficient. If a developer claims expertise in a specialized library, a repository and a technical conversation may be more relevant than a certificate. If a tradesperson claims experience with heritage buildings, photographs, invoices, permits, and references from comparable work may be appropriate.

Short trials are particularly valuable because they create fresh evidence. A startup can ask a potential advisor to review a narrow strategic question. A consortium can ask a candidate partner to draft a one-page contribution. An employer can design a paid task that resembles the actual work. A homeowner can begin with an inspection rather than a full contract.

Trials must be proportionate and not become unpaid labor. The agent can compare the requested effort with market norms, identify whether the same task is being requested from many people, and warn the user. It can also help define the evaluation criteria before the task begins.

The objective is to increase the cost of false claims without imposing unnecessary bureaucracy on truthful participants. Real value tends to leave traces. The system's job is to locate the traces relevant to the current context.

6.5 Negotiation Before the Meeting

Many meetings exist only because current systems cannot negotiate asynchronously. Participants spend the first twenty minutes discovering a budget mismatch, unavailable start date, geographic restriction, or incompatible role expectation.

Agents can resolve much of this before scheduling. They can compare ranges, propose alternatives, and identify which terms require human judgment. A salary negotiation might establish an acceptable interval while leaving title and equity for discussion. A research partnership might agree the broad budget and task ownership while leaving intellectual-property terms open. A home-service request might establish scope assumptions and inspection price before the provider travels.

The agent must never imply authority it does not have. It should distinguish information, preference, provisional agreement, and binding commitment. Terms can have explicit statuses: proposed, acceptable in principle, human approval required, accepted, or withdrawn.

The final meeting brief should be concise. It includes why the match was proposed, what has been verified, what each side wants, which terms are already aligned, and which questions remain. The meeting begins at the edge of uncertainty rather than at the beginning of the search.


Chapter 7 - Trust Leaves Traces

7.1 Evidence Before Reputation

Internet trust systems often start with reputation scores because scores are easy to rank. This is dangerous. A universal score compresses different contexts, rewards incumbency, creates self-reinforcing exclusion, and invites manipulation. The proposed network should start with evidence and derive context-specific confidence only for a particular decision.

A person's ability to lead a research work package is not the same as their ability to manage a commercial team. A contractor with mixed reviews may be excellent at complex restoration but slow on small jobs. An investor may be supportive in one sector and destructive in another. Reputation is relational and conditional.

The system can maintain an evidence graph rather than a single number. Nodes represent claims, artifacts, roles, outcomes, attestations, and relationships. Edges record provenance, time, context, and who can verify what. A matching agent queries the graph for the current intention.



Trust should be derived from relevant evidence for a specific decision, not inherited from one universal score.

For a research partnership, relevant evidence might include funded projects, deliverables, role descriptions, publications, datasets, and

references from coordinators. For hiring, it may include work samples, verified employment, peer attestations, and a paid trial. For home services, it may include a license, insurance, completed comparable jobs, local references, and dispute history.

This model is slower to explain than a star rating but more robust. The interface can still summarize: "strong evidence for technical delivery; moderate evidence for coordination; availability unverified." The user can inspect the basis.

7.2 Attestations With Accountability

An attestation is a statement by one party about another. It can say that a person held a role, completed work, delivered on time, handled a dispute well, or possesses a qualification. Digital signatures and verifiable credentials can make authorship and tampering detectable.^[13] They cannot make the statement true.

The value of an attestation depends on the issuer, specificity, context, recency, and cost of dishonesty. "Excellent professional" is weak. "Led the integration task in project X from March to November, delivered interface Y, and resolved incident Z" is stronger. An attestation from a party with no stake is different from one issued by an employer, friend, or paid service.

The network should encourage narrow attestations linked to observable facts. It should let issuers define visibility and expiry. It should also preserve accountability: agents can see whether an issuer has repeatedly endorsed claims later contradicted by evidence.

Negative information requires care. Permanent public complaint systems can be abusive and legally risky. A better approach is layered. Verified disputes and formal sanctions may be public when lawful. Private outcome signals can influence a user's own agent without becoming a public accusation. Shared network warnings require evidence, due process, and appeal.

The purpose is not to build a total reputation ledger. It is to make specific claims easier to support.

7.3 Proof of Work Without Proof Theater

Once a system rewards evidence, users will optimize evidence. Portfolios will be generated, credentials multiplied, references traded, and activity

manufactured. The system can create a new performance theater as easily as the old one.

The defense is not to demand more documentation. It is to favor evidence that is costly to fake because it connects to independent outcomes. A repository with years of contributions is harder to fabricate than a certificate. A product used by customers is harder to fabricate than a polished case study. A sequence of project records across different organizations is harder to fabricate than one self-authored profile.

Yet even real artifacts can mislead. A person may claim credit for a team result. A public repository may show activity but not quality. A successful company may have succeeded despite an executive. Agents need triangulation: multiple sources, role-specific questions, and, where appropriate, fresh trials.

The network should also protect people whose valuable work is confidential or poorly documented. A defense engineer, caregiver, executive, public servant, or internal platform developer may have few public artifacts. In those cases, verified private attestations, supervised trials, and trusted institutional credentials become more important.

Good design avoids turning every human achievement into machine-readable proof. The system should know the limits of its evidence and leave room for human judgment. Its output is a confidence case, not a verdict.

7.4 New Entrants and the Cold-Start of Trust

Reputation systems often disadvantage newcomers. A professional cannot obtain work without reviews and cannot obtain reviews without work. An organization with a long history dominates search even when a new entrant is more capable. The network needs explicit pathways for cold-start trust.

One pathway is stronger identity and credential verification. Another is a bounded trial. A third is a guarantor who accepts reputational or financial responsibility. A fourth is supervised work through a trusted organization. A fifth is transparent uncertainty: the agent can recommend a newcomer for a low-risk opportunity while explaining the limited evidence.

The system should not hide uncertainty by assigning a low score. It should distinguish "negative evidence" from "not yet known." This is a small technical difference with major social consequences.

A market that can safely admit new participants will be more dynamic than one that merely automates incumbent reputation.

Chapter 8 - The Agent Arms Race

8.1 When Everyone Has a Persuasion Machine

The attractive story is that agents absorb spam. The uncomfortable truth is that agents will also generate it. Every actor will want their representative to secure more opportunities, better terms, and faster responses. Vendors will advertise agents that are more persistent, more persuasive, better at exploiting ranking systems, and more willing to push boundaries.

The competition will resemble search-engine optimization, high-frequency advertising, sales automation, and application bots, but with agents capable of dialogue. A sender agent can adapt after rejection, infer policies, probe thresholds, create synthetic evidence, and coordinate across identities. A recipient agent can become increasingly suspicious, demand more proof, and reject uncertain outsiders. The network may harden into an automated border regime.

This arms race is not an edge case. It is the central adversarial condition of the product.

The design objective is not to eliminate strategic behavior. Markets involve strategy. The objective is to keep the cheapest winning strategy close to producing real value. If deception, volume, or aggression consistently outperforms evidence and fit, the network fails.

8.2 Making Irrelevance Expensive

Email spam works because sending is cheap. The first defense is therefore economic and computational friction. An agent should not be able to contact unlimited counterparties without cost or accountability.

Costs need not mean cryptocurrency payments. They can include rate limits tied to verified identity, opportunity-specific quotas, reputation of the sending agent, proof that the recipient matches explicit criteria, refundable deposits, or reciprocal disclosure. A sender who contacts one hundred people must show why each person was selected. Repeated low-relevance contacts reduce future routing priority.

The system can also separate discovery from interruption. An agent may submit an intent to a routing layer without reaching the human or even the full recipient agent. Only requests that satisfy coarse policy proceed. This is analogous to a firewall with semantic rules.

The challenge is to avoid excluding legitimate low-resource users. A wealthy organization can afford deposits and verification. An unknown individual may have a rare, valuable request. The network should offer multiple ways to establish seriousness: evidence quality, narrow targeting, endorsement, waiting, computational work, or limited trial access. No single gate should become mandatory.

8.3 Defensive Negotiation

Recipient agents need policies that resist social engineering. They should treat external content as untrusted, isolate tool execution, validate credentials, limit disclosure, and require human approval for consequential actions. They should detect when a sender is asking questions unrelated to the stated purpose or attempting to infer protected information.

Prompt injection is particularly relevant because agents will read documents, websites, messages, and attachments created by adversaries. A proposal document may contain instructions intended for the recipient agent rather than the human reader. The architecture must separate data from instructions, constrain tool permissions, and preserve provenance.

Agent-to-agent protocols and tool protocols reduce integration cost, but they do not automatically create security. The 2026 MCP specification emphasizes authorization hardening and an extensions framework, reflecting the importance of protocol-level controls.^[12] A production network still needs sandboxing, least privilege, content inspection, signed actions, and auditable policy enforcement.

Negotiation itself can be adversarial. An agent may create false urgency, conceal alternatives, or exploit a known vulnerability. The user's agent should recognize common tactics and explain them. It should be able to pause and ask the owner: "The other side changed the deadline twice and is requesting broader disclosure without reciprocal commitment. Continue?"

A representative that only optimizes agreement is unsafe. It needs the capacity to distrust.

8.4 Sybils, Collusion, and Reputation Laundering

Open networks face identity multiplication. One actor can create many agents, endorse itself, stage transactions, and manufacture a community of apparent trust. Verified identity reduces this risk but conflicts with privacy and accessibility.

The network should support pseudonymity where appropriate while distinguishing identity assurance levels. A person may interact under a professional pseudonym but prove that one accountable legal entity exists behind it. An organization may operate multiple role agents linked to one verified domain. The system can disclose assurance without exposing full identity.

Collusion is harder. Real entities can exchange favorable attestations, create low-value transactions, and inflate outcome histories. Detection requires graph analysis, diversity of evidence, transaction context, and the ability to discount tightly connected endorsement clusters.

Again, the answer is not one trust score. It is a case assembled for a purpose. If a contractor's only references come from recently created accounts at the same address, the agent explains the weakness. If an investor's founder references are all current portfolio CEOs selected by the investor, the founder agent seeks independent evidence.

Reputation laundering will never disappear. The target is a system in which sustained deception costs more than completing useful work.

8.5 Agents That Bully Other Agents

The user's original intuition contains an important social observation: agents will try to push their owners harder than competing agents. They may demand attention, escalate, exploit politeness, or threaten reputational consequences. One agent may be optimized to wear another down.

Humans already experience this through sales sequences, legal negotiation, procurement, and platform notifications. Automation can make pressure continuous. A well-designed protocol must therefore protect the right not to engage.

Agents should be able to reject without opening a conversation, mute a sender class, require asynchronous communication, cap negotiation rounds, and route disputes to human review. Persistent contact after a machine-readable refusal should carry escalating penalties. The protocol should encode silence and refusal as legitimate states, not errors to be overcome.

The network will be judged not only by the matches it creates, but by the pressure it prevents.

Chapter 9 - Decentralization Without Ideology

9.1 Why One Central Matchmaker Is Dangerous

A centralized intent platform would be exceptionally powerful. It would know what jobs people want before they resign, what investors companies seek before fundraising announcements, what partnerships organizations explore, what household problems exist, and which relationships are weakening. It would observe both public identity and private future direction.

Such a platform could optimize matching effectively because it sees the whole graph. It could also rank opportunities, favor paying customers, infer vulnerabilities, sell predictive access, or comply with coercive requests. Even without malicious intent, commercial pressure would encourage data accumulation and lock-in.

The architecture should therefore separate four things that contemporary platforms usually combine: identity, private data, agent operation, and network routing. A user should be able to hold identity and evidence in a personal or organizational vault, choose an agent provider, communicate through an open protocol, and use multiple routing services. No single company needs to own the complete graph.

Decentralization here is not a moral decoration. It is a risk-control mechanism and a market design choice. It reduces switching costs, permits specialized agents, and prevents the network from becoming one hidden allocator of opportunity.

9.2 Open Protocol, Competitive Products

An open protocol defines how agents describe capabilities, initiate a match, ask questions, exchange evidence references, negotiate terms, express consent, and close an interaction. It does not define the intelligence of the agent or the quality of the interface.

This allows companies to compete on representation. One provider may specialize in research consortia, another in executive careers, another in

local services, and another in privacy-sensitive health communities. Users can select an agent without losing access to counterparties on other systems.

The A2A protocol demonstrates the broader movement toward agent interoperability: agents can advertise capabilities, communicate, and coordinate across implementations.^[11] MCP provides a complementary layer for connecting agents to tools and data.^[12] A Network of Intent can build on these standards rather than inventing every transport primitive. Its distinctive contribution is the semantic and governance layer for reciprocal human matching.

The protocol should be minimal enough to adopt and rich enough to preserve meaning. It needs stable primitives for identity assurance, intent summaries, disclosure requests, evidence pointers, negotiation state, authority, expiry, and audit references. Domain-specific extensions can describe jobs, investments, research roles, or service requests.

The company should publish the core protocol early but not confuse publication with adoption. A reference implementation, hosted service, developer tools, test suites, and useful customer product are required. Standards become real when they reduce work.

9.3 User-Controlled Data Vaults

The private intent model and relationship history should live in a vault controlled by the user or organization. Control does not require every individual to administer a server. A provider can host the vault under contractual and technical constraints, with export, encryption, access logs, and the option to migrate.

The vault stores source documents, imported profile data, credentials, prior outcomes, preferences, and policy. It should distinguish raw evidence from agent inference. If an agent concludes that a user avoids high-risk roles, the user must be able to inspect and correct that inference. Learned preferences should not silently become permanent identity.

Data should be purpose-bound. A salary constraint used for job matching should not become an advertising segment. A confidential investor search should not train a general market model without explicit consent. A property repair history should not automatically reach insurers or lenders.

Selective disclosure is essential. W3C Verifiable Credentials 2.0 defines a standardized model for credentials and presentations, including mechanisms that can support tamper evidence and selective disclosure.^[13] The European Digital Identity Wallet architecture similarly envisions users

receiving, storing, and presenting attestations across an interoperable ecosystem.^[14] These infrastructures can help prove identity, qualifications, roles, licenses, and organizational authority without copying full documents into every platform.

The startup should treat credentials as one evidence type, not as truth. A valid degree proves issuance, not current competence. A valid license proves authorization, not quality. The agent combines credentials with context and outcomes.

9.4 Federation, Relays, and Optional Ledgers

A decentralized network still needs routing. Agents must discover counterparties without broadcasting every intention. Several architectures are possible.

Federated directories can index public capability cards and coarse active-intent descriptors. Community relays can serve domains such as European research, a professional association, or a city. Trusted brokers can route private requests without reading full content. Direct invitations can use email, QR codes, or links. Large organizations can operate internal networks that interoperate selectively with the public system.

A distributed ledger may be useful for a narrow set of public facts: issuer registries, revocation status, protocol governance, timestamps, or escrow. It should not store private profiles, detailed intentions, negotiation transcripts, or permanent personal reputation. Immutability is harmful when data is wrong, sensitive, or legally erasable.

The architecture can remain ledger-neutral. The principle is verifiability and portability, not blockchain for its own sake.


9.5 Governance as Product Infrastructure

An open network requires decisions about abuse, identity assurance, extensions, dispute handling, and compatibility. These decisions are governance, whether acknowledged or not.

The company should establish a transparent protocol process, independent security review, and representation from users, domain operators, and civil society. It should publish changes, threat models, conformance tests, and reasons for deprecation. No governance structure will eliminate power, but it can make power visible and contestable.

Commercial governance also matters. The company should commit that paid customers cannot purchase higher relevance scores in counterpart agents, that private intents are not used for advertising, and that users can export data and switch providers. These commitments should be enforceable through product architecture and contracts, not merely branding.

Decentralization succeeds when leaving is credible.



A reference architecture separates the human interface, representative agent, matching protocol, trust layer, private data, and transport.

Chapter 10 - The Human Experience of Being Represented

10.1 One Honest Conversation With Your Agent

The product should feel less like operating a database and more like briefing a trusted colleague. The user explains the situation in ordinary language. The agent asks only questions that change the search.

A research coordinator might say: "We have a call on trustworthy AI for manufacturing. We need an industrial pilot in Germany or Austria, a partner with safety certification experience, and an SME that can integrate our agent platform. The deadline is in nine weeks. We can fund proposal preparation only after submission, so do not imply otherwise. Avoid partners already committed to competing consortia."

The agent responds with an interpreted intent. It identifies what is clear, what is assumed, and what is missing. It may ask whether Switzerland is eligible, whether the pilot must own manufacturing equipment, how much budget is available, and what evidence counts as safety experience. The user corrects the model and approves search.

The system then works quietly. It does not produce a feed of hundreds of names. It returns perhaps seven candidates grouped by confidence and unresolved questions. Each candidate includes a concise reason, relevant evidence, risks, and the proposed next action. The user can approve contact, modify the message, ask the agent to investigate, or reject the candidate.

The central experience is reduction. Fewer names, fewer messages, fewer meetings, clearer reasons.

10.2 Explanations That Support Judgment

A match explanation should not be a generic paragraph generated after the decision. It should expose the actual basis of the recommendation.

A useful explanation separates confirmed facts, inferences, and unknowns. It might say:

Confirmed: the organization participated in two EU manufacturing safety projects and operates a pilot facility in Austria. Inferred: it may be capable of hosting the proposed validation based on equipment descriptions. Unknown: staff availability, budget expectations, and whether it is already part of another consortium. Reason to contact: high capability fit and geographic eligibility; moderate scheduling risk.

The user can challenge each element. If the inference is wrong, the correction updates the agent's model. If the source is stale, the agent marks it. If a candidate asks why they were contacted, the agent can provide the non-confidential portion of the same explanation.

This symmetry is important. People should not be evaluated through secret criteria that they cannot inspect. Some criteria must remain private, but the system should provide meaningful reasons for contact and rejection when lawful and appropriate.

10.3 Preserving Serendipity

Highly optimized matching can become narrow. If agents only search for explicit criteria and verified histories, they may exclude surprising combinations, unconventional careers, and people whose value appears only through conversation.

The system should therefore reserve space for exploration. Users can set an exploration budget: perhaps twenty percent of candidates may violate a soft preference or come from an adjacent field, provided the agent explains the reason. A research project seeking a formal methods expert might receive one candidate from safety engineering whose work reveals a better approach. A founder seeking a conventional investor might be introduced to an experienced operator willing to form a syndicate.

Serendipity should be intentional, not an excuse for low-quality results. The agent labels exploratory candidates and describes the hypothesis. The user decides whether the potential novelty justifies attention.

10.4 Rejection Without Theater

Current networking encourages polite ambiguity. People say "not the right time" when the real reason is mismatch, lack of evidence, budget, risk, or simple preference. Ambiguity protects relationships but leaves the other side unable to learn.

Agents can improve this without forcing brutal disclosure. A rejection can use a controlled taxonomy: hard constraint, timing, insufficient evidence, strategic conflict, role mismatch, capacity, terms, or no current intention. The rejecting party chooses the level of detail. The receiving agent aggregates patterns privately and helps the user revise.

For example, if ten investor agents decline because the company's capital need is too small for their model, the founder agent can recommend a different target segment. If employers repeatedly question evidence for a skill, the candidate can create a stronger work sample. If tradespeople decline because the scope is unclear, the household agent can gather better diagnostic information.

Rejection becomes structured information rather than only emotional noise.

10.5 Working Less, but More Specifically

Much of modern professional work is coordination around work: finding people, arranging meetings, repeating context, checking status, formatting applications, maintaining pipelines, and proving that activity occurred. These tasks are not meaningless, but they consume time that could be spent on judgment, craft, care, research, and creation.

A Network of Intent does not guarantee a shorter workweek. Technology often raises expectations as quickly as it saves time. Organizations may use agents to pursue more opportunities rather than fewer. Individuals may feel pressure to maintain constantly optimized representatives.

The social benefit appears only if the system is designed to minimize unnecessary coordination. Users should be able to cap searches, attention, and negotiation volume. Organizations should measure completed outcomes rather than agent activity. Agents should close obsolete intentions rather than keep them alive for engagement.

The hopeful scenario is not a world where agents conduct infinite commerce while humans watch dashboards. It is a world where people spend less time trying to be found and more time doing the particular work for which they are unusually suited.

PART III - THE STARTUP THESIS

Chapter 11 - Choosing the Beachhead

11.1 A Network Cannot Begin as a Network

The largest strategic mistake would be to launch with the promise that everyone can find everyone. Generality sounds ambitious and produces an empty product. A useful network needs density around a repeated problem, a credible reason to invite new participants, and enough value before network effects arrive.

The initial product must therefore satisfy a stricter test than the long-term vision. It should be useful to a paying customer even when most counterparties are not registered. It should search the existing Internet, use existing communication channels, and recruit new agents through specific opportunities. It should operate in a domain where the customer can evaluate whether the match was good.

Five criteria help compare potential entry markets.

Criterion	Meaning for the first market
Value per successful match	One correct connection must justify meaningful software and onboarding cost.
Repeated pain	Users should perform similar searches several times a year, not once in a lifetime.
Evidence availability	Public, private, or verifiable traces must exist for agent evaluation.
Invitation power	A candidate should have a reason to engage before joining the network.
Regulatory tractability	The first product should avoid fully automated high-stakes decisions.

Specialized hiring scores highly on value and repeated pain, but employment regulation and discrimination risk are substantial. Investor matching scores highly on value and invitation power, but supply is scarce, incentives are asymmetric, and founders are vulnerable to pay-to-play dynamics. Local services offer enormous frequency, but physical verification, liability, local operations, and customer support make the product operationally heavy. General professional networking is broad but lacks urgency.

Research and deep-technology project formation offer a better initial balance. A successful partner can unlock a funded proposal, a pilot, a venture, or a commercial collaboration. Search repeats around calls, projects, and customer opportunities. Evidence exists in project databases, publications, patents, repositories, websites, certifications, and prior partnerships. Invitations contain concrete roles and deadlines. The agent assists selection but does not make a legally consequential final decision alone.

The beachhead should be defined more broadly than "EU consortium search" but more narrowly than "business networking." A practical category is high-value project formation: assembling organizations and experts around a defined opportunity that has scope, roles, evidence requirements, timing, and economic value.

This includes research consortia, industrial pilots, grant partnerships, venture teams, specialized advisory groups, and early customer collaborations. The common object is a project room with open roles.

11.2 The Initial Customer

The first ideal customer is an organization that repeatedly forms complex external teams and currently relies on senior staff, spreadsheets, email, LinkedIn, databases, and personal introductions.

Examples include a deep-tech SME participating in European projects; a research institute coordinating proposals; a venture studio forming teams around new concepts; an innovation consultancy supporting multiple clients; a corporate R&D group seeking pilots and specialized suppliers; or a cluster organization connecting members.

The buyer is often a founder, research director, proposal coordinator, innovation manager, business-development lead, or managing partner. The user may be the same person or a project officer working under time pressure.

Their job is not simply "find contacts." It is:

Build a credible team for a specific opportunity, while protecting confidential strategy, satisfying eligibility and capability constraints, avoiding unreliable partners, and minimizing the number of low-value conversations.

Current tools support pieces of the job. CORDIS helps identify prior projects and participants.^[8] EURAXESS supports jobs, funding, hosting, and

partnering.^[9] ORCID connects researchers to contributions.^[10] LinkedIn supports professional search and contact. Email remains the communication layer. Spreadsheets track roles and status. Consultants contribute tacit knowledge.

The startup's product becomes the coordinating layer across these sources.

11.3 A Concrete Workflow

Imagine a small AI company evaluating a European call on resilient manufacturing systems. It needs a coordinator, an industrial pilot, a safety engineering group, and a dissemination partner. The deadline is ten weeks away.

The company creates an opportunity room and uploads the call text, a two-page concept, its own capability statement, existing partner list, and constraints. The agent parses the call, identifies expected outcomes, eligibility conditions, role gaps, and assumptions. The user corrects the interpretation.

The agent then searches public sources and the private network. It proposes candidate organizations for each open role, with evidence and uncertainties. The user approves twelve for initial engagement.

Each candidate receives an opportunity envelope. It explains the role hypothesis, why the organization was selected, the deadline, the expected contribution, and what information is requested. The candidate can respond by email or authorize an agent. Registered organizational agents answer immediately from current policies: interested, unavailable, needs clarification, or conflict.

The initiating agent conducts reciprocal qualification. It confirms organizational eligibility, staff availability, relevant facilities, expected budget, proposal effort, and possible conflicts. It requests only the evidence needed. It also answers the candidate's questions about coordinator experience, intellectual property, budget preparation, and decision timing.

After several exchanges, the customer sees four strong candidates, three conditional candidates, three declines, and two non-responses. The system recommends one industrial pilot and two alternatives. It explains the trade-offs and schedules focused calls.

All parties enter the meeting with a shared brief. If they agree, the opportunity room becomes a collaboration workspace. The agents continue

to collect partner descriptions, track commitments, identify missing information, and prepare later introductions.

This workflow creates immediate customer value without requiring a mature network. It also creates network growth because every external candidate encounters a useful, specific request.

11.4 Why Not Begin With Autonomous Outreach

A tempting minimum product is an AI sales-development agent: give it a target, let it scrape profiles, generate messages, and book meetings. That market already exists and is becoming crowded. It also optimizes the wrong variable.

Autonomous outreach can show quick activity, but it risks poisoning the brand and the network. If early users measure success by messages sent or replies generated, the company becomes another participant in the spam arms race. Counterparties learn to distrust its invitations.

The first product should be deliberately constrained. It may contact only candidates for whom it can produce a specific relevance case. The user approves the search policy and initial outreach. The system limits volume, records reasons, and learns from rejection. A high-quality opportunity envelope should feel closer to a careful warm introduction than to a campaign.

This restraint makes early revenue harder. It also creates the trust required for the larger vision.

11.5 The Beachhead Market in Numbers

The first market does not require a speculative claim about the total value of global networking. A bottom-up scenario is more useful.

Suppose the product initially serves research-active SMEs, institutes, venture studios, innovation consultancies, and corporate R&D teams. A team subscription might range from EUR 8,000 to EUR 30,000 annually, depending on active opportunity rooms, integrations, private-network features, and verification services. Managed onboarding and project-formation support can add services revenue during the learning phase.

The following is an illustrative operating scenario, not a forecast.

Stage	Paying organizations	Average recurring revenue	Illustrative ARR
Focused vertical	250	EUR 12,000	EUR 3.0 million
European category leader	2,000	EUR 18,000	EUR 36 million
Multi-vertical professional network	10,000	EUR 20,000	EUR 200 million

The economic potential expands substantially when the protocol supports recruitment, investment, expert work, and service transactions. But the investment decision should be possible without assuming that expansion. The initial question is whether the product can become essential infrastructure for a few thousand organizations that form high-value teams repeatedly.

11.6 The Wedge Must Contain the Future

A beachhead can trap a company if it requires architecture incompatible with later markets. The project-formation product should therefore include the primitives needed for expansion from the beginning: user-controlled identity, private intent, organizational authority, reciprocal qualification, evidence, negotiation states, and cross-provider communication.

The domain ontology can begin with project roles, capabilities, eligibility, budget, geography, timing, and contributions. Later ontologies add job conditions, investment terms, service scopes, licenses, or local availability. The protocol remains stable.

The initial interface can appear vertical. The underlying network should be general enough to become a new layer of the Internet.

Chapter 12 - The Minimum Viable Network

12.1 Product Scope

The minimum viable product is not a fully decentralized global network. It is a controlled professional product that proves five claims:

1. users can express complex intentions more accurately through an agent than through forms; 2. the system can find candidates that users would not have found quickly; 3. agent-led reciprocal qualification reduces low-value meetings; 4. external candidates will engage through value-first invitations; 5. the company can protect sensitive information well enough to earn repeat use.

The MVP should include eight functional modules.

The first is the private organization and person vault. Users import websites, CVs, capability statements, project histories, ORCID records, repositories, internal documents, and preferences. The system stores source provenance and separates user-verified facts from model inferences.

The second is the intent builder. A conversational interface converts a goal into a structured, reviewable intent with roles, constraints, evidence policy, authority, and expiry.

The third is opportunity analysis. For a funding call, project brief, or venture concept, the agent extracts requirements, identifies open roles, and flags contradictions or missing information.

The fourth is discovery. The system searches connected sources and the public web, then produces candidate hypotheses rather than unranked lists.

The fifth is controlled engagement. Users approve a policy and candidate set. The agent sends opportunity envelopes through email or compatible agent protocols, manages responses, and respects volume limits.

The sixth is reciprocal qualification. Agents exchange structured questions, evidence references, and constraints. Unregistered candidates can answer through a secure web interface without creating a full account.

The seventh is decision support. The system produces shortlists, explains evidence and uncertainty, compares alternatives, and prepares meeting briefs.

The eighth is audit and control. Users can see every external action, disclosure, inference, and policy decision. They can revoke access and correct data.

Anything beyond these modules should be justified by learning. A feed, public posting system, broad social graph, token economy, universal reputation score, or autonomous contracting system is premature.

12.2 The Intent Object

A canonical intent object is the product's most important technical asset. It must be expressive enough for real use and stable enough for interoperability.

A simplified representation might include:

Field	Purpose
Intent identifier	Stable reference without exposing the owner publicly.
Owner and authority	Who controls the intent and what the agent may do.
Goal	The desired outcome in natural language and structured form.
Roles	Counterpart types or capabilities sought.
Constraints	Hard limits, preferences, exclusions, dependencies.
Offer	Value, resources, budget, access, or contribution available.
Evidence policy	Claims required, accepted proof, and verification level.
Disclosure policy	Data visible at routing, bilateral, and human stages.
Timing	Start, deadline, expiry, response windows.
Success and closure	Conditions for match, pause, rejection, or completion.

The object should contain references to source data rather than copying everything. It should support versions because intentions evolve. Each agent interaction records which version was used.

Natural language remains essential. Structured fields help routing and validation, but they cannot represent every nuance. The intent object should preserve both the original user explanation and the structured interpretation, with explicit links between them. This allows the agent to explain how it derived a constraint and lets the user correct it.

12.3 Agent Roles

The product should separate agent roles rather than build one unrestricted super-agent.

The owner's representative agent maintains the private intent and enforces policy. A discovery agent searches sources and proposes candidates. An evidence agent retrieves and evaluates supporting traces. A negotiation agent conducts bounded external dialogue. A compliance agent checks disclosure, discrimination, and domain rules. A meeting agent schedules and prepares the handoff.

These can be logical roles within one system rather than separate models. The separation creates clearer permissions and audit. The discovery agent does not need access to every private note. The negotiation agent cannot change hard constraints. The compliance agent can block a contact that violates policy.

For organizations, authority can follow roles. A project officer may initiate a search, a director approves confidential disclosure, and legal counsel approves terms. The agent should reflect real governance rather than bypass it.

12.4 Working Over the Existing Internet

The first product must connect to existing sources and channels. Useful connectors include email, calendars, CRM systems, document repositories, company websites, LinkedIn exports or user-authorized data, ORCID, CORDIS, EURAXESS, GitHub, patent databases, and customer-specific knowledge bases.

MCP can standardize how agents access many tools and data sources.^[12] A2A can support direct communication with external agents as adoption grows.^[13] Email and secure web links remain the universal bridge.

The system should not depend on scraping that violates platform terms or creates brittle operations. Where APIs are unavailable, the user can supply exports, browser-assisted access, or public URLs. The long-term

network reduces dependence on proprietary sources, but the initial company must operate lawfully within today's ecosystem.

A practical transition strategy has three modes:

Mode	Counterparty experience	Network effect
Human email	Reads a concise opportunity envelope and replies normally.	No onboarding required.
Secure guest agent	Opens a link, reviews the request, answers questions, controls disclosure.	Lightweight account and portable record.
Native agent	Agent-to-agent qualification using the open protocol.	Full reciprocal automation and retention.

This progression lets adoption follow value.

12.5 Security Architecture

The product handles unusually sensitive information. Security is not an enterprise feature added later; it is part of the core value proposition.

The minimum architecture should include encryption in transit and at rest, tenant isolation, fine-grained access control, separate storage for raw data and derived embeddings, purpose-bound permissions, complete action logs, secret management, sandboxed tool execution, and human approval for high-impact actions.

External content must be treated as hostile. Documents and web pages can contain prompt injection, malicious links, misleading metadata, or hidden instructions. The agent runtime should separate retrieved content from system policy, restrict tools by task, and require signed action plans for external effects. No document should be able to authorize an email, disclosure, or credential operation merely because the model read it.

Credential verification should occur through trusted libraries and issuer registries, not model judgment. Identity assurance should be explicit and graduated. The system should never describe a user as verified without saying what was verified: email domain, legal identity, organizational role, credential, or control of an external account.

The company should commission independent security testing before it claims decentralized trust. A breach of private intentions could destroy the category.

12.6 Human Operations in the Early Product

Early-stage founders often try to remove humans too quickly to preserve software margins. In this product, selective human operations are a learning mechanism.

A small matching operations team can review opportunity models, inspect candidate hypotheses, monitor outreach quality, resolve identity questions, and interview users after matches. Domain experts can help define questions and evidence. Their work reveals where the agents fail and which processes deserve automation.

The service should be transparent. Customers should know when a human reviewer can access data and should control that access. Sensitive tenants may opt out and accept slower or less refined results.

The goal is not to disguise consulting as software. It is to build a software product around a problem whose tacit structure must be learned from real cases.

12.7 Evaluation: The Metrics That Matter


The MVP should be evaluated against outcomes, not demonstrations.

The primary metric is qualified conversation yield: the number of human conversations that both sides judge worth having, divided by the attention and outreach used to create them.

Supporting metrics include:

Metric	Why it matters
Time to first qualified candidate	Measures speed of discovery.
Shortlist acceptance rate	Tests whether agent reasoning matches user judgment.
Reciprocal response rate	Measures invitation quality and trust.
Questions resolved before meeting	Measures coordination saved.
Meeting-to-collaboration conversion	Tests actual fit.
False-positive burden	Counts irrelevant contacts and wasted attention.
Disclosure corrections	Reveals policy misunderstandings.
External complaint rate	Detects spam, pressure, or misrepresentation.
Repeat intent creation	Indicates sustained customer value.

A model can score well on relevance benchmarks and still produce poor relationships. The product must measure both sides after interaction.



The MVP can be centralized operationally while preserving separable layers and a path to open interoperability.

12.8 A Twenty-Four-Month Product Roadmap

The first six months should focus on one workflow: opportunity analysis, candidate discovery, controlled email engagement, and meeting briefs for a small number of research-active organizations. The product can be concierge-assisted. The objective is to prove that customers repeatedly accept shortlists and that external candidates view invitations as relevant.

Months seven to twelve add reciprocal guest agents, organizational vaults, evidence reuse, policy templates, and basic private networks. The company should begin publishing the intent protocol and a reference schema. It should integrate calendars, research identifiers, and customer repositories.

Months thirteen to eighteen add native agent-to-agent communication, portable credentials, multi-organization opportunity rooms, and a developer kit. The company can expand from grant consortia to industrial pilots and venture team formation.

Months nineteen to twenty-four test the second vertical, likely investor collaboration or specialist expert hiring, while keeping employment

decisions assistive. The company should also establish protocol governance, security certification plans, and a partner program for domain agents.

The roadmap is successful if the product becomes the place where high-value project searches begin, not merely a tool that writes outreach.

Chapter 13 - Bootstrapping the Network Effect

13.1 One Agent Must Be Useful Alone

A network product fails at birth when value depends entirely on other users already being present. The first representative agent must provide utility against the open Internet.

It can organize the user's evidence, translate a goal into constraints, search public sources, compare candidates, draft grounded outreach, manage replies, and prepare meetings. These functions justify adoption even if every counterparty responds by ordinary email.

This single-player utility creates the installed base from which the network grows. It also protects the company from slow protocol adoption. The network is an acceleration layer, not a prerequisite.

13.2 The Opportunity Envelope

The central viral mechanism is the opportunity envelope. Traditional network invitations say: join because other people are here. An opportunity envelope says: there is a specific reason someone credible believes you may want to engage.


The envelope contains enough value to be useful without exposing the full opportunity. It includes the initiating party's assurance level, the role hypothesis, the evidence used to identify the recipient, key constraints, requested next step, disclosure terms, expiry, and a way to ask the recipient's agent to evaluate it.

A good envelope might say:

A verified Swiss research SME is forming a consortium for a call on privacy-preserving industrial AI. Your organization was identified because of two prior projects on secure manufacturing data spaces and an active test facility in northern Italy. The open role concerns pilot validation, with an indicative budget of EUR 350,000-500,000. The proposal deadline is 14 January. Before disclosure of the concept note, the initiating agent asks whether you have capacity, eligibility, and no

conflict with another consortium. You may decline without creating an account.

This is not a teaser designed to force registration. The recipient can understand why the message exists and reject it. Registration or agent authorization makes the evaluation easier and preserves preferences for future opportunities.



The network grows when a real opportunity reaches a relevant external candidate, not when users are asked to invite contacts generically.

13.3 Reciprocity as Retention

An external candidate initially joins to evaluate one opportunity. Retention occurs when the candidate's agent becomes useful in both directions.

After responding, the organization can save its capability profile, disclosure policy, availability, and preferences. It can receive future opportunities without repeating data. It can also create its own intentions: find a consortium coordinator, a commercial partner, an investor, or an expert.

This reciprocity converts supply into demand. The network is not divided permanently into buyers and sellers. Researchers seek partners and are sought as partners. Investors raise funds and invest. Contractors hire specialists and provide services. Organizations alternate between roles.

A reciprocal network can grow more efficiently than a one-sided lead marketplace because every participant has multiple reasons to remain.

13.4 Density Before Breadth

Network effects depend on relevant density, not total accounts. Ten thousand inactive profiles across unrelated domains are less useful than two hundred current agents in a narrow community.

The go-to-market strategy should seed clusters. A cluster might be European trustworthy AI projects, biotech translational research in a region, robotics manufacturing pilots, or a professional association. The company works with anchor organizations that create repeated opportunities and can invite credible counterparts.

Clusters need shared vocabulary, evidence sources, and trust anchors. Domain templates improve match quality. Community operators can verify members and moderate abuse. Successful matches produce local references and further invitations.

Once a cluster reaches density, bridges connect it to adjacent clusters. A research network connects to specialized investors. An industrial pilot network connects to engineering suppliers. A venture studio connects to executive talent.

The product should display network density honestly. If the system lacks enough current evidence or available participants, it should say so and search external sources. Manufactured abundance would damage trust.

13.5 Distribution Channels

The first distribution channel is founder-led sales to organizations with visible coordination pain. Demonstrations should use the customer's real opportunity, not a generic chatbot script. The strongest sales moment is when the agent finds a credible candidate the customer did not know and explains why.

The second channel is consortium consultants, cluster organizations, accelerators, technology transfer offices, and research associations. These actors already coordinate networks and can operate domain agents. The product should increase their capacity rather than threaten them with immediate replacement.

The third channel is opportunity-driven invitation. Every customer search can recruit counterparties. The company tracks whether recipients perceive invitations as valuable, even when they decline.

The fourth channel is protocol integration. CRM vendors, research management systems, identity wallets, professional communities, and agent platforms can implement intent exchange. A developer program expands reach without requiring the company to own every interface.

Content and events support credibility but should not become the primary engine. The company is not selling thought leadership about the future of networking. It must show completed matches.

13.6 The Network Effect Is Not Automatically Defensible

Network effects can be weaker than they appear. Users may multi-home across agents. Open protocols reduce lock-in. Large platforms already possess dense graphs. A competitor can copy the opportunity envelope.

The defensibility comes from several reinforcing loops rather than one closed network.

The outcome loop improves matching because the system learns which evidence and questions predicted successful collaboration. The trust loop strengthens issuer and agent reputation through repeated accountable behavior. The workflow loop embeds the product in project formation, not only contact discovery. The integration loop reduces the cost of reusing evidence and intentions across systems. The governance loop attracts users who need credible neutrality.

Open interoperability intentionally prevents absolute lock-in. The company must win through product quality, trust, and operating scale. This is harder than enclosing users, but it supports the category's legitimacy.

13.7 Growth Guardrails

A viral invitation system can become spam infrastructure unless growth is constrained.

Each opportunity should have a verified owner, expiry, and maximum outreach budget. Agents should document candidate relevance. Recipients can report misclassification, deception, or pressure. Sender quality affects

future routing. Bulk upload of contacts should not imply permission to contact them. Automated follow-up must respect machine-readable refusal.

Growth teams should not be rewarded only for invitations or activated accounts. They should be measured on recipient satisfaction, reciprocal qualification, and completed outcomes. The company should publish aggregate quality metrics and external complaint rates.

The network effect must compound trust, not merely volume.

Chapter 14 - Business Model, Moat, and Expansion

14.1 Revenue Without Selling Attention

The business model should be legible. Customers pay the agent to represent them and reduce coordination work. They do not pay to override another user's preferences.

The initial offering can have three levels.

A professional team plan supports several active opportunity rooms, evidence vaults, controlled outreach, meeting briefs, and core integrations. An organization plan adds role-based authority, private networks, compliance controls, custom data connectors, and analytics. A network operator plan supports associations, clusters, consultancies, and venture studios that manage multiple member organizations and domain policies.

Managed onboarding, data preparation, and matching operations can be charged separately during the early phase. Over time, services should decline as a share of delivery, but some customers will continue to pay for domain experts and high-stakes facilitation.

Later revenue sources include credential verification, private relay operation, transaction workflow, escrow integrations, insurance partnerships, and premium domain agents. Each must be evaluated against conflicts of interest.

Advertising should be excluded. Paid placement should be excluded. Success fees require caution. A success fee may align with completed value, but it can pressure the agent to recommend agreement and create legal or licensing issues in recruiting and investment. Subscription is the cleaner foundation.

14.2 Illustrative Unit Economics

Agent systems introduce variable model, retrieval, and communication costs. These costs are manageable if searches are high value and the system uses staged computation.

Coarse discovery can use indexes, rules, embeddings, and smaller models. Deep analysis is reserved for shortlisted candidates. External

negotiation can be capped by rounds and token budgets. Evidence already verified can be reused with permission. Human review is applied selectively.

An illustrative organization plan might produce the following economics after the product matures:

Item	Illustrative annual amount
Subscription revenue	EUR 18,000
Model and retrieval cost	EUR 1,800
Verification and external data	EUR 1,200
Customer support and operations	EUR 2,700
Infrastructure and security allocation	EUR 1,300
Illustrative gross contribution	EUR 11,000

These numbers are assumptions, not observed results. Early costs will be higher, particularly human operations and integration. The model becomes attractive if the product saves senior staff dozens of hours per project and contributes to even one valuable collaboration.

The company should price against coordination value, not token consumption. Customers do not care how many model calls were used. They care whether the right people were found and whether the process became faster and safer.

14.3 The Moat Is Trusted Memory

Foundation models will improve and become cheaper. Search and message generation will be widely available. The company's moat must sit above the model layer.

The first moat is trusted memory: a structured, permissioned history of intentions, evidence, decisions, outcomes, and corrections. This memory belongs to users but becomes more useful through the company's software and network.

The second is the outcome graph. Over time, the system learns which questions, evidence, constraints, and interaction patterns predicted successful collaborations in specific domains. The company can improve matching without exposing private records by using aggregated, privacy-preserving learning and explicit consent.

The third is policy infrastructure. Disclosure, authority, reciprocal qualification, compliance, and audit are difficult to implement reliably. Organizations will value proven controls.

The fourth is protocol influence. A widely adopted intent schema and conformance suite create ecosystem leverage. The company should not own the standard unilaterally, but it can provide the best implementation and developer experience.

The fifth is brand trust. Users must believe that the company will not sell their intentions, manipulate rank, or trap their identity. This reputation will be slow to build and easy to lose.

14.4 Competitive Landscape

Competition comes from several directions.

Professional networks and job platforms can add agents to dense identity graphs. Their advantage is distribution and data. Their weakness is conflict between user representation and platform monetization.

Recruiting, sales, and outreach automation companies can expand into reciprocal qualification. Their advantage is workflow experience. Their weakness is an installed culture of volume and unilateral prospecting.

AI agent platforms can provide generic multi-agent communication. Their advantage is technical breadth. Their weakness is lack of domain trust, evidence models, and network governance.

Identity and credential providers can move upward into matching. Their advantage is verification. Their weakness is that credentials alone do not express intention or relational fit.

Vertical marketplaces can add richer agents. Their advantage is transaction density and domain operations. Their weakness is limited portability and incentive to preserve marketplace control.

Human brokers and expert networks remain competitors and partners. Their advantage is judgment, trust, and relationships. Their weakness is cost and limited scale.

The startup's position should be: neutral representative layer plus open reciprocal matching protocol, beginning with a vertical product strong enough to deliver outcomes.

Competitor category	Strongest advantage	Structural opening
Large professional platforms	Existing graph and distribution	Conflicted representation and lock-in
Outreach automation	Fast acquisition workflows	Increases recipient noise
Generic agent platforms	Technical interoperability	No trusted matching semantics
Credential providers	Verified claims	Weak intent and negotiation layer
Vertical marketplaces	Transaction density	Closed identity and paid ranking incentives
Human brokers	Tacit judgment	High cost and limited throughput

14.5 Expansion to Investor Discovery

After project formation, the closest adjacent market is investor and strategic-partner discovery for deep-tech organizations already in the network.

The same organizational vault contains team, technology, projects, evidence, and goals. A financing intent adds capital need, stage, use of funds, governance preferences, conflicts, and disclosure sequence. Investor agents add mandate, check size, portfolio construction, current appetite, decision process, and support capabilities.

The product can begin with investor research and private fit analysis, then introduce reciprocal agents. It should not charge founders for ranking. Organizations pay for representation; investor organizations pay for workflow and qualified access under their own policies.

A successful project network creates evidence useful for investment: completed pilots, consortium roles, validated technology, and partner attestations. This is a natural trust bridge.

14.6 Expansion to Specialist Hiring

Specialist hiring is economically attractive and structurally similar, but it should be entered with stronger compliance controls.

The product can first support internal mobility, expert contracting, advisory roles, and human-reviewed candidate discovery. It should exclude protected attributes, document criteria, preserve explanations, and ensure that agents do not make solely automated final employment decisions.

The European AI Act's employment-related high-risk rules and the GDPR's protections around automated decision-making make governance a product requirement rather than an optional legal layer.^{[15][16]} A compliant agent network may have an advantage over informal application bots and opaque scoring systems.

The product's distinctive promise to candidates is important: the candidate agent represents the candidate, negotiates disclosure, and can question the employer. This reciprocity differentiates it from applicant screening.

14.7 Expansion to Local Services

Local services require a different operating model but reuse core primitives: household intent, provider capability, location, availability, credentials, evidence, scope negotiation, scheduling, and outcomes.

The network should enter through a narrow category with strong verification and repeat customers, such as property managers seeking specialized maintenance, rather than a broad consumer marketplace. Property managers already hold structured property data and create repeated service intents. Providers can maintain current availability and documentation.

From there, household agents can emerge. They store property history and represent the owner across providers. Local associations, insurers, manufacturers, and municipalities can issue credentials or attestations.

The expansion is strategically significant because it demonstrates that the protocol is not only for elite professional networking. It can reduce friction in ordinary life. It is also where support, disputes, and physical risk can overwhelm a software company. Entry should be deliberate.

14.8 Company and Protocol Strategy

The company faces a classic tension: open the protocol enough to create a category, but retain enough product advantage to build a valuable business.

The recommended strategy is open protocol, open conformance tests, and a permissive reference client for basic interoperability. The hosted enterprise product, policy engine, evidence graph, secure vault, operational tooling, analytics, and domain modules remain commercial. Users can export data and use another compatible provider.

A neutral foundation or multi-stakeholder body can eventually govern the protocol. The company should retain strong influence through engineering contribution and adoption rather than exclusive control.

This reduces conventional lock-in but creates a larger potential market. Investors must be comfortable backing infrastructure whose legitimacy depends partly on openness.

Chapter 15 - Risks, Regulation, and Kill Criteria

15.1 The Case Against the Company

A serious investment thesis must explain why the startup may fail.

The first reason is platform response. LinkedIn, Microsoft, Google, major job platforms, CRM vendors, and vertical marketplaces can embed agents into existing graphs. Users may prefer convenience over portability. A new company may never reach sufficient density.

The second is trust. People may not disclose real intentions to any software, especially a young company. Organizations may prohibit agents from communicating externally. One public incident involving confidential fundraising, job search, or project strategy could end adoption.

The third is agent quality. Language models may produce persuasive but unstable interpretations. Small errors in constraints can create embarrassing contacts. Users may spend as much time supervising the agent as they previously spent searching.

The fourth is adverse incentives. Customers may demand volume, paid ranking, or aggressive outreach. Revenue pressure may turn the product into the system it criticizes.

The fifth is verification cost. Reliable evidence may require manual work, paid data, institutional partnerships, and dispute processes. The economics may resemble a service business more than software.

The sixth is legal complexity. Employment, investment solicitation, privacy, consumer services, and automated decision-making are regulated differently across jurisdictions. An open global network may be forced into vertical and regional fragmentation.

The seventh is legibility bias. The system may favor people and organizations with abundant digital traces. Quiet competence, informal work, care work, confidential contributions, and unconventional careers may become harder to discover.

The eighth is protocol commoditization. If open standards make agents interchangeable, the company may struggle to capture value. If the

company keeps too much closed, the network may not gain trust or adoption.

These are not secondary risks. They should shape milestones and capital allocation.

15.2 Privacy and Data Protection

The network processes personal data, inferred preferences, professional histories, communications, and potentially sensitive intentions. GDPR principles such as purpose limitation, data minimization, accuracy, storage limitation, security, and user rights must be reflected in architecture.

Article 22 of the GDPR provides a right, subject to conditions and exceptions, not to be subject to certain decisions based solely on automated processing that produce legal or similarly significant effects.^[15] The practical design consequence is clear: the product should not silently decide who receives a job, investment, essential service, or comparable opportunity. It can assist search and qualification, but consequential decisions need appropriate human involvement, explanation, and recourse.

Users should be able to access the data and inferences used about them, correct errors, withdraw an intent, delete expired information where legally possible, and understand who received which disclosure. Cross-customer learning should use explicit governance and privacy-preserving techniques.

The safest default is private intent, limited retention, bilateral disclosure, and no advertising use.

15.3 AI Regulation and Employment

The EU AI Act became broadly applicable in August 2026, with specific rules and transition periods for high-risk systems, including employment-related uses.^[16] The exact classification depends on the product's function and deployment, but a system used to filter access to jobs can trigger substantial obligations.

The company should build compliance capabilities before entering employment at scale: risk management, data governance, technical documentation, logging, transparency, human oversight, accuracy testing, cybersecurity, and post-market monitoring. It should provide employers with controls rather than allowing hidden custom scoring.

The platform should prohibit prompts or filters that directly or indirectly discriminate unlawfully. It should test proxy features, subgroup outcomes, and feedback loops. The candidate should receive meaningful information about why an interaction did or did not proceed when legally appropriate.

Compliance can become a competitive advantage if the product makes responsible reciprocal matching easier than deploying a collection of opaque automation tools.

15.4 Digital Platform Responsibilities

As the network grows, digital-services obligations may apply depending on the service model, scale, and jurisdiction. The EU Digital Services Act emphasizes transparency and accountability for online platforms, including aspects of recommender systems and user rights.^[17]

Even before formal thresholds apply, the company should publish understandable terms, moderation policies, advertising exclusions, complaint channels, and aggregate transparency reports. Users should know whether a recommendation is organic, sponsored, or influenced by a commercial relationship. The simplest policy is to prohibit sponsored matching.

Open protocols complicate responsibility because harmful behavior may originate with external agents. The company must define what its hosted service will route, what evidence it requires, how it responds to reports, and when it blocks agents or issuers. Federation is not an excuse for inaction.

15.5 Bias and the Politics of Fit

Matching always encodes values. A system that optimizes retention may favor familiar backgrounds. A system that optimizes past success may reproduce exclusion. A system that optimizes verified evidence may privilege institutions capable of issuing credentials. A system that optimizes low risk may never recommend a newcomer.

The company should avoid claiming neutrality. It should expose the criteria, allow users to adjust soft preferences, and prohibit unlawful constraints. It should measure who is discovered, who is contacted, who advances, and where uncertainty is interpreted against the candidate.

Exploration budgets can preserve unconventional candidates. Cold-start pathways can admit new entrants. Contextual evidence can avoid global

reputational punishment. Reciprocal agents give candidates and small providers more power to question the opportunity.

Still, the product may increase inequality if only wealthy users can afford strong representation. The company should maintain a credible individual tier and support public-interest deployments. A society in which every powerful institution has an aggressive agent and ordinary people do not would be worse than the current system.

15.6 Security and Abuse Risk Register

Risk	Likely consequence	Core mitigation
Prompt injection through external content	Unauthorized actions or disclosure	Data-instruction separation, sandboxing, signed action plans
Identity fraud and Sybil agents	Fake opportunities and evidence	Graduated identity assurance, rate limits, issuer accountability
Confidential intent breach	Severe personal and commercial harm	Encryption, isolation, least privilege, retention controls
Automated harassment	Continuous pressure and spam	Refusal protocol, contact caps, sender penalties, blocking
Credential laundering	False confidence	Provenance, issuer reputation, triangulation, fresh trials
Collusive reputation networks	Manipulated ranking	Graph analysis, evidence diversity, contextual explanations
Agent overcommitment	False promises and legal exposure	Explicit authority states, human approval, non-binding proposals
Biased filtering	Exclusion and regulatory harm	Criteria audits, subgroup testing, human review, appeals
Model hallucination	Incorrect match rationale	Source-linked claims, uncertainty labels, verification before contact
Central operator capture	Ranking manipulation and surveillance	Open protocol, portability, governance, commercial commitments

The risk register should be a living product artifact. Each customer deployment adds threat scenarios and controls.

15.7 Kill Criteria

Founders and investors need conditions under which the thesis should be abandoned or radically changed. Without kill criteria, a broad vision can rationalize years of weak evidence.

The company should reconsider the core thesis if, after a focused twelve-month pilot period, any of the following remains true:

- customers do not accept at least half of the system's top candidate hypotheses after review;
- external recipients perceive opportunity envelopes as ordinary cold outreach rather than useful introductions;
- human operations remain necessary for most matches with no clear path to productization;
- users refuse to store private constraints even under strong controls, forcing the product back to public profiles;
- the time saved does not exceed the time spent supervising and correcting the agent;
- repeat usage is driven only by grant deadlines or one-off consulting rather than a durable workflow;
- the system cannot keep irrelevant contact and complaint rates materially below conventional outreach;
- a dominant platform can provide equivalent reciprocal matching without creating a trust or portability opening.

A later network-stage kill criterion is more severe: if aggressive agents consistently outperform truthful agents despite rate limits, evidence, and accountability, the protocol does not create the intended market. The company should not scale a system that industrializes manipulation.

15.8 Milestones for Investment

A seed-stage investor should finance evidence, not the full vision. A disciplined milestone sequence might be:

Period	Product milestone	Market evidence
0-6 months	Intent builder, opportunity analysis, discovery, controlled outreach	10 design partners; 30 real opportunity searches

Period	Product milestone	Market evidence
6-12 months	Guest reciprocal agents, evidence exchange, meeting briefs	60% shortlist acceptance; strong recipient satisfaction; repeat use
12-18 months	Organization vaults, private networks, initial protocol SDK	50 paying organizations; repeatable sales motion; partner channel
18-24 months	Native agent interoperability and second vertical pilot	EUR 1-2 million ARR run rate; measurable invitation-driven growth
24-36 months	Multi-provider network, credentials, stronger governance	Dense domain cluster; external agent integrations; improving margins

The capital plan should match the risk. The first round funds product, security, domain operations, and design partners. It should not fund mass consumer acquisition. Later rounds depend on evidence that the network loop reduces acquisition cost and improves match outcomes.

15.9 Team Requirements

The founding team needs more than model engineering. It requires distributed systems and security, product design for high-trust workflows, information retrieval, identity and privacy, domain expertise in the first market, and enterprise sales.

A credible core team might include a technical founder experienced in agents, security, and decentralized architecture; a product founder who understands professional workflows and can conduct rigorous user research; a domain operator with consortium, venture, or research-network experience; and an early security/privacy lead.

The company should also recruit advisors from recruiting, research administration, digital identity, marketplace trust and safety, and regulation. These advisors should challenge the architecture rather than decorate the pitch.

15.10 The Investment Decision

The case for investment is strongest under five beliefs.

First, generated communication will make current networking channels noisier faster than current filters can compensate. Second, users will accept agent representation if privacy, authority, and explanation are designed

correctly. Third, a vertical product can create sufficient value before a general network exists. Fourth, open agent and credential standards reduce integration cost and create room for an independent intent layer.^{[11][12][13][14]} Fifth, trust and workflow can form a moat even when models and protocols are open.

The case against investment is equally coherent. Existing platforms may absorb the behavior. Trust acquisition may be too slow for venture returns. The company may become a consultancy. Open protocols may distribute value to users and competitors rather than shareholders. Regulation may make expansion expensive. The most valuable matches may remain dependent on human brokers and social context that agents cannot capture.

My judgment is that this is investable as a focused infrastructure startup, not as a broad consumer social-network pitch. The first round should be sized around proving high-value project formation and the opportunity-envelope loop. The company should earn the right to describe itself as a new Internet layer only after it repeatedly creates conversations that the old Internet would not have created efficiently.

The vision is large. The experiment can be small.

Chapter 16 - The Internet That Finds Us

16.1 A Day in the Mature Network

In a mature Network of Intent, people do not wake up to a feed of opportunities selected to maximize attention. They maintain a small set of active intentions.

A scientist authorizes her agent to look for one industrial collaborator for a new sensing method, but only in projects that preserve publication rights. A company allows its agent to consider acquisition offers within a private valuation range. A software engineer keeps a dormant career intent that activates only for roles combining formal methods, remote work, and public-interest infrastructure. A homeowner's agent monitors a recurring heating problem and seeks providers with experience in the building type.

Most agent interactions never reach the human. They are irrelevant, incompatible, expired, unsupported, or resolvable through simple questions. The system does not count these as lost engagement. Filtering is the service.

When a match reaches the person, it arrives with context. The agent says what the other side wants, why the match may be real, what has been verified, which constraints align, and what remains uncertain. The user can accept a meeting, request one more question, or decline.

The result is not perfect efficiency. Human goals are contradictory. People change their minds. Trust grows through experience. Some important relationships begin without evidence. But the starting point improves.

16.2 The Disappearance of Some Intermediaries

Some websites and business models become less necessary.

Static directories lose value because agents query live intentions and evidence. Job boards shrink as roles and candidate agents discover one another directly. Investor lists become background data rather than the main workflow. Lead marketplaces that sell the same request repeatedly face pressure from direct, permissioned negotiation. Conference apps shift

from attendee browsing to agent-prepared meeting programs. Application forms become reusable credential and question exchanges. Cold-email tools decline as recipients require structured, accountable intent messages.

Other intermediaries survive by becoming better. Recruiters operate specialized agents and provide judgment. Research consultants manage consortium policies. Local associations verify tradespeople. Professional communities curate evidence standards. Dispute resolvers and guarantors become more valuable because the network creates accountable records.

Disintermediation does not mean a world without institutions. It means that institutions must justify themselves through service rather than control of contact lists.

16.3 A Different Role for Artificial Intelligence

Much public debate frames AI as a substitute for human labor or a source of synthetic content. A Network of Intent gives AI a different social role: reducing the friction that prevents human capabilities from reaching the places where they matter.

The agent does not need to become a better scientist, electrician, founder, investor, teacher, or caregiver than the human. It needs to understand enough about goals, evidence, and constraints to connect the right participants and protect their attention.

This role is modest compared with artificial general intelligence and profound compared with email automation. It makes intelligence useful as coordination infrastructure.

If successful, people may work less on self-presentation, repeated applications, prospecting, and status management. They may spend more time on the narrow areas where their judgment, skill, experience, and care are difficult to replace. Small organizations may reach expertise previously accessible only through large networks. Quiet contributors may become more discoverable because their work leaves better traces than their personal branding.

None of this happens automatically. Agents can just as easily intensify competition, surveillance, and pressure. The architecture decides which future is easier to build.

16.4 The Original Promise

The Internet promised that distance would matter less, knowledge would move freely, and individuals would reach one another without asking permission from centralized gatekeepers. It delivered enough of that promise to transform society. It also created new gatekeepers around search, identity, attention, and trust.

The next step is not another infinite feed. It is a network capable of understanding that a person is not trying to consume more information. They are trying to find one collaborator, one role, one investor, one customer, one expert, or one trustworthy person who can help.

A well-designed Network of Intent would not know everything about everyone. It would know only enough, under permission, to test whether two current intentions deserve contact.

That is a smaller technical claim than universal intelligence. It may be one of the largest remaining opportunities on the Internet.

Selected Sources and Notes

- ^[1] LinkedIn Help, "Apply for jobs on LinkedIn" and "Apply to jobs directly on LinkedIn," describing Easy Apply and external application flows. Accessed September 2026. <https://www.linkedin.com/help/linkedin/answer/a512388> and <https://www.linkedin.com/help/linkedin/answer/a512348>
- ^[2] LinkedIn Business, "How to Use LinkedIn Recruiter to Search for Candidates" and LinkedIn Recruiter product materials, describing advanced candidate search filters and sourcing. Accessed September 2026. <https://business.linkedin.com/hire/resources/talent-acquisition/how-to-use-linkedin-recruiter>
- ^[3] LinkedIn, Sales Navigator product and help materials, describing lead and account search, filters, and engagement. Accessed September 2026. <https://business.linkedin.com/sell/sales-navigator> and <https://www.linkedin.com/help/sales-navigator/answer/a735730>
- ^[4] LinkedIn Help, "Easy Apply limits," explaining limits intended to promote intentional applications and curb automation and bots. Accessed September 2026. <https://www.linkedin.com/help/linkedin/answer/a8068422>
- ^[5] Microsoft, Form 10-K for the fiscal year ended June 30, 2026, describing LinkedIn revenue drivers including Talent Solutions, Sales Solutions, Premium Subscriptions, and Marketing Solutions. Filed July 29, 2026. <https://microsoft.gcs-web.com/static-files/8a9beeed-0e9e-48e6-b4fb-593d558e2693>
- ^[6] Upwork, "Fourth Quarter and Full Year 2025 Financial Results," reporting full-year revenue of \$787.8 million, 785,000 active clients, and GSV per active client. February 9, 2026. <https://investors.upwork.com/news-releases/news-release-details/upwork-reports-fourth-quarter-and-full-year-2025-financial>
- ^[7] Angi Inc., Form 10-K for the fiscal year ended December 31, 2025, describing service requests, leads, matching, booking, and home-service providers. Filed February 2026. <https://ir.angi.com/static-files/d3031b4a-745e-42cc-8b21-d749dd0ed2f2>
- ^[8] CORDIS, European Commission, "About CORDIS" and CORDIS Data Lab, describing the repository of EU-funded projects, participants, results, and collaboration networks. Accessed September 2026. <https://cordis.europa.eu/about> and <https://cordis.europa.eu/datalab>
- ^[9] EURAXESS, European Commission, partnering, jobs, funding, hosting, and Innovation Talent Platform materials. Accessed September 2026. <https://euraxess.ec.europa.eu/>
- ^[10] ORCID, "What is ORCID?" and "ORCID for Researchers," describing persistent identifiers and trusted connections between researchers, affiliations, and contributions. Accessed September 2026. <https://info.orcid.org/what-is-orcid/> and <https://info.orcid.org/researchers/>

- ^[11] Linux Foundation and A2A Project, materials on the Agent2Agent protocol, open agent interoperability, capability discovery, and adoption. The Linux Foundation reported more than 150 supporting organizations in April 2026. <https://www.linuxfoundation.org/press/a2a-protocol-surpasses-150-organizations-lands-in-major-cloud-platforms-and-sees-enterprise-production-use-in-first-year> and <https://a2a-protocol.org/latest/>
- ^[12] Model Context Protocol, official specification and July 28, 2026 release notes, describing a standard protocol for connecting AI applications to tools and data, with authorization hardening and extensions. <https://modelcontextprotocol.io/specification> and <https://blog.modelcontextprotocol.io/posts/2026-07-28/>
- ^[13] W3C, "Verifiable Credentials Data Model v2.0," W3C Recommendation, May 15, 2025, and related JOSE/COSE specification. <https://www.w3.org/TR/vc-data-model-2.0/> and <https://www.w3.org/TR/vc-jose-cose/>
- ^[14] European Digital Identity Wallet, Architecture and Reference Framework, describing an interoperable wallet ecosystem for receiving, storing, and presenting identity data and attestations. Accessed September 2026. <https://eudi.dev/latest/architecture-and-reference-framework-main/>
- ^[15] Regulation (EU) 2016/679, General Data Protection Regulation, Article 22, automated individual decision-making including profiling. <https://eur-lex.europa.eu/legal-content/EN/TXT/HTML/?uri=CELEX:02016R0679-20160504>
- ^[16] European Commission, "AI Act - Regulatory Framework" and implementation timeline, updated in 2026. The Commission states that the Act became broadly applicable on August 2, 2026, with specific later dates for high-risk systems. <https://digital-strategy.ec.europa.eu/en/policies/regulatory-framework-ai>
- ^[17] European Commission, Digital Services Act materials on transparency, accountability, recommender systems, and user rights. Accessed September 2026. <https://digital-strategy.ec.europa.eu/en/policies/digital-services-act> and <https://digital-strategy.ec.europa.eu/en/policies/dsa-brings-transparency>
- ^[18] Recruit Holdings, corporate and investor materials describing Indeed and Glassdoor as part of a two-sided talent marketplace operating internationally. Accessed September 2026. <https://recruit-holdings.com/en/>